

INSTITUTO TECNOLÓGICO DE AERONÁUTICA



Eduardo Moura Zindani

**EDUCATIONAL ORBIT SIMULATION WITH
GENERATIVE AI AGENTIC WORKFLOW AND
VIRTUAL REALITY VISUALISATION**

Final Paper
2025

Course of Aerospace Engineering

CDU ????

Eduardo Moura Zindani

**EDUCATIONAL ORBIT SIMULATION WITH
GENERATIVE AI AGENTIC WORKFLOW AND
VIRTUAL REALITY VISUALISATION**

Advisor

Prof. Dr. Christopher Shneider Cerqueira (ITA)

AEROSPACE ENGINEERING

SÃO JOSÉ DOS CAMPOS
INSTITUTO TECNOLÓGICO DE AERONÁUTICA

2025

Cataloging-in Publication Data
Documentation and Information Division

Zindani, Eduardo Moura
Educational Orbit Simulation with Generative AI Agentic Workflow and Virtual Reality
Visualisation / Eduardo Moura Zindani.
São José dos Campos, 2025.
114p.

Final paper (Undergraduation study) – Course of Aerospace Engineering– Instituto Tecnológico de Aeronáutica, 2025. Advisor: Prof. Dr. Christopher Shneider Cerqueira.

1. AI. 2. VR. 3. Orbit. I. Instituto Tecnológico de Aeronáutica. II. Educational Orbit Simulation with Generative AI Agentic Workflow and Virtual Reality Visualisation.

BIBLIOGRAPHIC REFERENCE

ZINDANI, Eduardo Moura. **Educational Orbit Simulation with Generative AI Agentic Workflow and Virtual Reality Visualisation**. 2025. 114p. Final paper (Undergraduation study) – Instituto Tecnológico de Aeronáutica, São José dos Campos.

CESSION OF RIGHTS

AUTHOR'S NAME: Eduardo Moura Zindani

PUBLICATION TITLE: Educational Orbit Simulation with Generative AI Agentic Workflow and Virtual Reality Visualisation.

PUBLICATION KIND/YEAR: Final paper (Undergraduation study) / 2025

It is granted to Instituto Tecnológico de Aeronáutica permission to reproduce copies of this final paper and to only loan or to sell copies for academic and scientific purposes. The author reserves other publication rights and no part of this final paper can be reproduced without the authorization of the author.

Eduardo Moura Zindani
Rua H8B, Ap. 235
12.228-461 – São José dos Campos–SP

EDUCATIONAL ORBIT SIMULATION WITH GENERATIVE AI AGENTIC WORKFLOW AND VIRTUAL REALITY VISUALISATION

This publication was accepted like Final Work of Undergraduation Study

Eduardo Moura Zindani

Author

Christopher Shneider Cerqueira (ITA)

Advisor

Prof. Dra. Maísa de Oliveria Terra
Course Coordinator of Aerospace Engineering

São José dos Campos: June 20, 2025.

Acknowledgments

...

*"If I have seen farther than others,
it is because I stood on the shoulders of giants."*

— SIR ISAAC NEWTON

Resumo

Métodos educacionais tradicionais frequentemente encontram dificuldades para transmitir conceitos complexos, espaciais e dinâmicos como os da mecânica orbital. A recente convergência entre a Realidade Mista (RM) de consumo e os sofisticados agentes de Inteligência Artificial (IA) Generativa apresenta uma oportunidade para criar um novo paradigma de interfaces de aprendizagem intuitivas e experenciais. Este trabalho detalha o projeto, desenvolvimento e demonstração de uma plataforma educacional interativa para a exploração dos princípios da mecânica orbital. O objetivo principal do sistema é conectar a teoria de leis físicas abstratas à compreensão intuitiva, permitindo que os usuários aprendam por meio da interação corporificada. A metodologia é centrada em uma arquitetura modular que integra dois componentes principais: (1) um "cérebro" agente generativo, impulsionado por Modelos de Linguagem Abrangentes, que interpreta comandos em linguagem natural e atua como um guia educacional especializado; e (2) um "mundo" de simulação em tempo real e visualização imersiva em realidade mista, construído no motor Unity para o Meta Quest 3, que renderiza trajetórias orbitais fisicamente precisas em um espaço tridimensional onde os usuários vivenciam a mecânica orbital de dentro. A plataforma facilita um ciclo de interação multimodal contínuo, onde os comandos de voz do usuário são capturados, processados pelo agente para alterar os parâmetros da simulação e refletidos na visualização imersiva em RV com feedback auditivo conversacional. Este trabalho entrega um protótipo funcional que demonstra uma nova abordagem para a educação científica, transformando dados abstratos em uma experiência manipulável e conversacional para promover uma aprendizagem exploratória e profundamente engajadora. A plataforma é disponibilizada como software de código aberto para permitir validação, adaptação e extensão pela comunidade para diversos contextos educacionais.

Abstract

Traditional educational methods often struggle to convey complex, spatial, and dynamic concepts such as those found in orbital mechanics. The recent convergence of consumer-grade Mixed Reality (MR) and sophisticated Generative AI agents presents an opportunity to create a new paradigm for intuitive and experiential learning interfaces. This paper details the design, development, and demonstration of an interactive educational platform for exploring the principles of orbital mechanics. The system's primary objective is to bridge the gap between abstract physical laws and intuitive comprehension by enabling users to learn through embodied interaction. The methodology is centered on a modular architecture that integrates two core components: (1) a generative agent "brain," powered by Large Language Models, which interprets natural language commands and acts as an expert educational guide; and (2) a real-time simulation and immersive mixed reality visualization "world," built in the Unity engine for the Meta Quest 3, which renders physically accurate orbital trajectories in a three-dimensional space where users experience orbital mechanics from within. The platform facilitates a seamless multimodal interaction loop where a user's voice commands are captured, processed by the agent to alter simulation parameters, and reflected in the immersive VR visualization with conversational auditory feedback. This work delivers a functional prototype that demonstrates a novel approach to science education, transforming abstract data into a manipulable, conversational experience to foster exploratory and deeply engaging learning. The platform is released as open-source software to enable community validation, adaptation, and extension for diverse educational contexts.

List of Figures

FIGURE 2.1 – End-to-end agentic workflow from Anthropic’s “Building Effective Agents.” The human issues a query through an <i>interface</i> ; the LLM asks clarifying questions until the task is precise, receives contextual files, iteratively writes and tests code against the environment, and finally returns results for display (Anthropic, 2024).	28
FIGURE 4.1 – Opening cutscene frame showing Earth from deep space. The camera begins at this distant perspective before zooming toward the Hub’s orbital position.	44
FIGURE 4.2 – Hub environment from user perspective. Earth floats in stereoscopic three-dimensional space, rendered with 8K textures to provide visual fidelity while maintaining 90 Hz frame rate on Quest 3.	45
FIGURE 4.3 – Close-up view of the simulation satellite in the Hub environment. The user has navigated their perspective to examine the spacecraft model.	47
FIGURE 4.4 – ISS Mission Space environment after scene transition. The pre-configured ISS orbit (420 km altitude, 51.6° inclination) is visible as a cyan trajectory. ANASTASIA, the ISS specialist, greets the user with context awareness, knowing they came to learn about altitude selection. This environment demonstrates the <code>MissionConfig</code> <code>ScriptableObject</code> system (Appendix A) that defines mission-specific orbits, characters, and knowledge domains.	51

FIGURE 4.5 – Close perspective view of the 422 km circular orbit after creation. The cyan trajectory traces a perfect circle around Earth’s equator (inclination 0°). The satellite model is visible at one point along the orbit. This visualization enables spatial observation of circular geometry that would be difficult to comprehend from 2D diagrams. The user can move their VR viewpoint to observe the orbit from multiple angles, supporting the spatial cognition discussed in Section 2.1.	56
FIGURE 4.6 – Hubble Mission Space with DR_HARRISON explaining elliptical orbital geometry. The environment shows Hubble’s circular orbit (540 km altitude, 28.5° inclination) as a reference for comparison.	59
FIGURE 4.7 – First elliptical orbit (periapsis 400 km, apoapsis 2,000 km) showing visible elongation compared to the circular orbit.	61
FIGURE 4.8 – Satellite at apoapsis (farthest point, slowest velocity) showing the ellipse’s oblong shape from a fixed perspective.	62
FIGURE 4.9 – Satellite at periapsis (closest point, fastest velocity) from the same perspective, demonstrating the geometric and kinematic contrast.	63
FIGURE 4.10 – Voyager Mission Space environment. CARL, the Voyager specialist, embodies a contemplative character distinct from the technical specialists encountered earlier. The voice synthesis uses slower pacing (speed 0.9, configured in <code>KarlVoiceSettings</code>) to convey thoughtful reflection. This environment demonstrates that the platform supports both experiential learning (Sections 4.2 and 4.3) and conceptual dialogue.	66

List of Tables

TABLE 4.1 – Tool Execution Summary from Demonstration Session. Note: Multiple visits to the same mission space (e.g., ISS visited twice) result in more scene transitions (9 total) than unique tool invocations shown here.	70
TABLE 4.2 – System Performance Metrics (Measured During Demonstration) . . .	71
TABLE 4.3 – Physics Validation: Simulation vs Published Mission Data	72
TABLE A.1 – Agent Prompt Component Specifications	82
TABLE A.2 – Tool Parameter Constraints	84
TABLE A.3 – Conversation Exchange Data Structure	85
TABLE B.1 – Physical Constants for Orbital Calculations	88
TABLE B.2 – Circular Orbit Calculation Algorithm	88
TABLE B.3 – Elliptical Orbit Parameter Constraints	90
TABLE B.4 – Scale Compression Examples	90
TABLE B.5 – LineRenderer Configuration Parameters	92
TABLE C.1 – Speech-to-Text Audio Capture Specifications	95
TABLE C.2 – Recording State Machine	95
TABLE C.3 – Text-to-Speech Synthesis Parameters	97
TABLE C.4 – Character Voice ID Assignments	98
TABLE C.5 – Voice Interaction Latency Budget	100
TABLE D.1 – Android Build Configuration	101
TABLE D.2 – Scene Build Index Configuration	102
TABLE D.3 – Controller Input Mapping	103

TABLE D.4 – VR Camera Hierarchy	104
TABLE D.5 – Spatial UI Rendering Configuration	105
TABLE D.6 – Common Scene Components	106
TABLE D.7 – Performance Targets for 90 Hz VR	107
TABLE D.8 – Frame Time Budget Breakdown (Hub Scene)	108

Contents

1	INTRODUCTION	18
1.1	Organisation	18
1.2	Motivation	19
1.3	Objectives	20
2	LITERATURE REVIEW	22
2.1	Augmented and Virtual Reality in Immersive Educational Simulation Systems	22
2.1.1	Hardware Evolution:	22
2.1.2	Software Ecosystems and Frameworks:	23
2.1.3	Use Cases in Education:	24
2.1.4	Embodiment, Interaction, and Spatial Cognition:	25
2.2	Generative Agents	26
2.3	Orbital Mechanics: The Physics of Celestial Motion	29
2.3.1	Newtonian Gravity and the Two-Body Problem	29
2.3.2	Kepler's Laws and Orbital Geometry	30
2.3.3	Orbital Regimes: Circular, Elliptical, and Hyperbolic Trajectories	31
2.3.4	Orbital Orientation: Inclination and Coverage	33
2.3.5	The Educational Foundation for Interactive Exploration	34
3	METHODOLOGY	36
3.1	Design Philosophy and Approach	36
3.2	System Architecture and Data Flow	37
3.2.1	Interaction Flow	37

3.3	Mixed Reality Design Rationale	38
3.4	Technical Implementation	38
3.5	Development and Version Control	38
3.6	Validation and Dissemination Strategy	39
3.7	Core Module Implementation	39
3.7.1	Agent System Implementation	39
3.7.2	Orbital Physics Simulation	40
3.7.3	Voice Integration Pipeline	41
3.7.4	Virtual Reality Environment	42
4	RESULTS AND DEMONSTRATION	43
4.1	Entering the Orbital Environment	44
4.1.1	Opening Sequence: Narrator Introduction	44
4.1.2	Hub Environment Arrival: Spatial Presence and Agent Welcome	45
4.1.3	Satellite Discovery: First Interaction	46
4.2	Learning Through Mission Specialist Dialogue: The ISS Circular Orbit	48
4.2.1	The Learner's Question: Seeking Conceptual Foundation	48
4.2.2	Scene Transition to ISS Mission Space: Modular Architecture in Action	50
4.2.3	Specialist Educational Dialogue: Real-World Context and Engineering Trade-offs	52
4.2.4	Returning to Hub and Creating the Circular Orbit: Applying Knowledge	53
4.2.5	Observing and Understanding Orbital Motion: Visual Physics and Time Control	55
4.3	Exploring Orbital Geometry: Elliptical Orbits and Eccentricity	58
4.3.1	Conceptual Question: Extending Beyond Circular Orbits	58
4.3.2	Hubble Specialist Explanation: Geometric and Operational Distinctions	59
4.3.3	First Elliptical Orbit Creation	60
4.3.4	Iterative Refinement	61
4.3.5	Visual Observation of Kepler's Second Law	62
4.4	Conceptual Extension: Escape Trajectories and Mission Context	64

4.4.1	Expanding Conceptual Boundaries: The Question of Deep-Space Missions	64
4.4.2	Voyager Specialist Dialogue: Bound vs Unbound Trajectories	65
4.5	System Integration and Technical Validation	69
4.5.1	Complete Tool Suite Execution and Usage Analysis	69
4.5.2	Real-Time Performance Metrics and System Latency	70
4.5.3	Physics Accuracy Validation Against Published Mission Data	72
4.5.4	Complete System Demonstration and Integration Evidence	74
4.6	Open-Source Platform Delivery	77
APPENDIX A – AGENT SYSTEM IMPLEMENTATION		82
A.1	Prompt Architecture	82
A.1.1	Hub Agent: Three-Tier Prompt System	82
A.1.2	Mission Specialist Prompts	83
A.2	Tool Schema and Validation	84
A.2.1	Tool Execution Pipeline	84
A.3	Conversation Context Management	84
A.3.1	Context Window Management	85
A.3.2	Cross-Scene Persistence	85
A.4	API Integration	85
A.4.1	Request Structure	85
A.4.2	Response Parsing	86
A.4.3	Mission-Specific Configuration	86
APPENDIX B – ORBITAL PHYSICS IMPLEMENTATION		87
B.1	Two-Body Keplerian Mechanics	87
B.2	Vis-Viva Equation Implementation	87
B.2.1	Circular Orbit Calculation	88
B.2.2	Elliptical Orbit Calculation	89
B.3	Scale Compression	90
B.3.1	Compression Factor Derivation	90

B.3.2	Numerical Stability	90
B.4	Trajectory Visualization	91
B.4.1	Orbital Ellipse Equation	91
B.4.2	Sampling Algorithm	91
B.4.3	Rendering Configuration	92
B.5	Coordinate System Conventions	92
APPENDIX C – VOICE PIPELINE IMPLEMENTATION		94
C.1	System Architecture	94
C.2	Speech-to-Text Pipeline	94
C.2.1	Push-to-Talk Input Detection	95
C.2.2	Audio Capture and Conversion	95
C.2.3	API Request Structure	96
C.2.4	Response Parsing	96
C.3	Text-to-Speech Pipeline	97
C.3.1	API Request Structure	97
C.3.2	MP3 Decoding and Playback	97
C.3.3	Model Selection Rationale	98
C.4	Character Voice Management	98
C.4.1	Scene-Specific Voice Switching	99
C.4.2	Voice Settings Persistence	99
C.5	Error Handling and Fallbacks	99
C.6	Performance Optimization	100
C.6.1	Memory Management	100
C.6.2	Latency Budget	100
APPENDIX D – VR DEPLOYMENT CONFIGURATION		101
D.1	Quest 3 Android Build Configuration	101
D.1.1	SDK Version Rationale	101
D.1.2	Stereo Rendering Pipeline	102
D.1.3	Build Index Scene Configuration	102

D.2 Input System Implementation	102
D.2.1 Push-to-Talk Implementation	102
D.2.2 Desktop Testing Mode	103
D.2.3 VR Mode Detection	103
D.3 Camera and Rendering Configuration	104
D.3.1 OVRCameraRig Structure	104
D.3.2 Near Clip Plane Configuration	104
D.3.3 Desktop Camera Alignment	104
D.4 Spatial UI Implementation	104
D.4.1 MissionClockUI Pattern	105
D.4.2 Transition Overlay System	105
D.5 Scene Architecture and Persistence	106
D.5.1 Singleton Persistence Mechanism	106
D.5.2 Asynchronous Scene Loading	107
D.6 Performance Optimization	107
D.6.1 Rendering Optimizations	107
D.6.2 Memory Management	108
D.6.3 Frame Time Breakdown	108
APPENDIX E – COMPLETE DEMONSTRATION TRANSCRIPT	109

1 Introduction

1.1 Organisation

This work is organised into three main chapters, each addressing a distinct aspect of the project. The breakdown is as follows:

- **Chapter 1: Introduction.** This chapter sets the stage for the research and development.
 - *Motivation* (§1.2): Presents the core argument that the convergence of immersive reality technologies—particularly Virtual Reality—and generative AI enables a new, more intuitive paradigm for educational interfaces.
 - *Objectives* (§1.3): Defines the project’s specific, actionable goals, centered on the development and demonstration of an interactive, agent-guided simulation platform.
- **Chapter 2: Literature Review.** This chapter provides the theoretical and technical foundation for the work by reviewing three key domains.
 - *Augmented and Virtual Reality* (§2.1): Reviews the evolution of immersive hardware and software ecosystems and establishes their pedagogical value for spatial learning.
 - *Generative Agents* (§2.2): Defines the architecture of modern LLM-powered agents, detailing their ability to use planning, memory, and external tools to reason through and execute complex tasks.
 - *Orbital Mechanics* (§2.3): Outlines the fundamental physics of celestial motion, including the two-body problem, classical orbital elements, and impulsive maneuvers, which form the mathematical basis for the simulation.
- **Chapter 3: Methodology.** This chapter details the practical design, implementation, and demonstration strategy of the project.

- *System Architecture and Data Flow* (§3.2): Describes the end-to-end fluxogram of the system, illustrating how user voice input is captured, processed by the agent, and rendered in the virtual reality simulation in a continuous loop.
- *Core Component Implementation* (§3.4): Details the specific development plan and tools for the two primary modules: the Generative Agent (“Brain”) and the Simulation and AR Visualisation (“World”).
- *Demonstration and Release Strategy* (§3.6): Outlines the approach for showcasing the platform’s capabilities through a demonstration video and releasing it as open-source software to enable community validation, adaptation, and extension.

1.2 Motivation

For decades, popular media and speculative fiction have envisioned futuristic interfaces for exploration and control, from holographic command centers to immersive planetary navigation tools. Films such as *Minority Report* (2002) and *Iron Man* (2008) popularized visions of humans interacting with vast information systems through gestures, speech, and spatial manipulation. These visions were once confined to science fiction, but today, the convergence of immersive reality technologies—particularly Virtual Reality (VR)—and Artificial Intelligence (AI) is bringing such interfaces into the realm of technological feasibility.

In particular, the past few years have seen rapid advances in consumer-grade immersive hardware, particularly in virtual reality. Devices like the Meta Quest and Apple Vision Pro represent significant milestones in accessibility and visual fidelity, enabling immersive environments that are no longer confined to laboratory research or elite applications. The implications for interface design, interaction paradigms, and knowledge acquisition are profound. VR is no longer a speculative technology; it is present, evolving, and increasingly democratized.

Concurrently, the emergence of generative AI and language-based agents has introduced a paradigm shift in how humans interact with complex systems. Large Language Models (LLMs), such as those powering conversational agents, can now interpret natural language, generate multimodal content, and coordinate sequences of actions across software environments. This represents a departure from deterministic, rule-based systems toward stochastic and adaptive workflows, where agents interpret intention, negotiate uncertainty, and build dynamically responsive experiences.

When these technologies—immersive VR and generative agents—are combined, they form the foundation for a new kind of interface: one that is spatial, conversational, and

adaptive. Such interfaces do not rely on code or static menus; they respond to voice, gesture, and embodied input. They transform abstract data into manipulable space, and procedural complexity into natural dialogue.

This is particularly relevant in the domain of education. Traditional educational systems remain bound to text, diagrams, and symbolic representation. While these tools are powerful, they often fall short when applied to fields that are inherently spatial, dynamic, or non-intuitive. Orbital mechanics, for example, involves motion through three-dimensional space governed by non-linear physical laws. Launch trajectories, gravitational slingshots, inclination changes, these are difficult to visualize and even harder to intuit.

In this context, immersive simulation becomes more than a visual aid: it becomes a cognitive bridge. A learner can speak a question and witness a launch trajectory materialize in their physical space. They can observe orbits evolve in real time, ask about inclinations or transfer windows, and receive explanations grounded in physics. Education becomes experiential, a process of exploration rather than instruction.

Moreover, generative agents provide a layer of accessibility that is historically absent in technical domains. They can guide the learner, interpret vague queries, correct misconceptions, and explain phenomena in adaptive ways. They act as intelligent mediators between curiosity and formal knowledge.

Given these technological conditions—the maturity of consumer VR, the rise of stochastic AI agents, and the persistent limitations of traditional educational media—this project is motivated by a clear opportunity: to construct a new type of educational experience. One that is not constrained by interface conventions, disciplinary jargon, or static presentation. One that invites the user to learn by seeing, asking, moving, and listening.

The convergence of embodied interaction and generative intelligence allows for a simulation system that is not only technically rigorous, but experientially meaningful. It enables a form of learning in which the abstract becomes tangible, the distant becomes near, and the user is placed at the center of the scientific process. This project emerges from the belief that space education, and scientific education more broadly, can and must evolve to meet the possibilities of our time.

1.3 Objectives

General Objective

To develop an interactive, agent-guided simulation platform that enables users to explore and understand orbital mechanics through embodied interaction, combining natural

language dialogue and real-time virtual reality visualizations.

Specific Objectives

1. Design and implement a simulation environment capable of rendering orbital trajectories in real time, grounded in physically accurate models.
2. Integrate a generative agent capable of interpreting natural language input, translating it into simulation parameters, and guiding the user through explanations and interactions.
3. Enable multimodal interaction by combining voice commands, spatial presence, and mixed reality visual feedback (VR immersion and AR passthrough capabilities) to create a seamless and intuitive user experience optimized for learning orbital mechanics concepts.
4. Ensure that all components of the system, simulation and agent, function coherently and communicate reliably in real time.
5. Create a system architecture that is modular and extensible, allowing for future expansion to other celestial bodies, educational modules, or mission types.
6. Demonstrate and deliver an open-source platform that showcases the potential for agent-guided virtual reality learning in orbital mechanics, enabling community validation, adaptation, and extension for diverse educational contexts.

2 Literature Review

2.1 Augmented and Virtual Reality in Immersive Educational Simulation Systems

Augmented Reality (AR) and Virtual Reality (VR) are complementary immersive technologies that enrich or replace a user's perception of the world. AR overlays digital content onto the real environment in real-time, allowing virtual objects to coexist with physical surroundings (Billinghurst; Clark; Lee, 2015). In contrast, VR completely immerses the user in a fully synthetic, computer-generated environment, blocking out the physical world. Milgram's classic "Reality-Virtuality" continuum illustrates these as end-points: AR lies near the real-world end (mixing virtual content with reality), whereas VR occupies the extreme virtual end with an entirely simulated world (Milgram; Kishino, 1994). In essence, AR adds to the user's real-world experience, while VR transposes the user into an interactive virtual scene. Both technologies share common roots in decades of research and development. The term augmented reality was first coined by Caudell and Mizell (1992) in the context of assisting Boeing manufacturing with see-through displays (Caudell; Mizell, 1992). A few years later, Azuma's influential survey defined AR by three key characteristics: combining real and virtual content, interactive operation in real time, and accurate 3D registration of virtual objects in the physical world (Azuma, 1997; Billinghurst; Clark; Lee, 2015). VR, meanwhile, has been long conceptualized as achieving presence – the feeling of “being there” in a virtual environment – by engaging multiple senses with responsive 3D graphics and audio (Johnson-Glenberg, 2018). Modern definitions emphasize that VR provides immersive first-person experiences where users can interact with simulated worlds as if they were real, inducing a strong sense of presence and agency within the virtual scene.

2.1.1 Hardware Evolution:

AR and VR technologies have evolved rapidly, enabling consumer-grade devices that support realistic immersive experiences. While early head-mounted displays date back

to the 1960s (e.g., Sutherland’s Sword of Damocles), the 2010s marked a turning point with modern devices. On the VR front, the Oculus Rift prototype (2010) by Palmer Luckey re-ignited interest with a wide field of view and affordable design. Crowdfunded in 2012 and acquired by Facebook in 2014, Oculus released its first consumer headset in 2016, alongside HTC’s Vive, which introduced room-scale tracking. These devices brought high-fidelity visuals and motion tracking to mainstream audiences.

The next major step came with standalone VR headsets. The Oculus/Meta Quest series, starting in 2019, integrated processing and inside-out tracking directly into the headset. Quest 2 (2020) and Quest 3 (2023) improved resolution, optics, and added passthrough AR capabilities (Ruth, 2024). In parallel, PC-based headsets like the Valve Index and Varjo pushed the fidelity frontier for gaming and enterprise simulation.

AR hardware followed a distinct trajectory. Initial systems used handheld or laptop setups, but the release of Microsoft’s HoloLens in 2016 marked the arrival of self-contained AR headsets with spatial mapping and inside-out tracking. Magic Leap One (2018) added novel display technologies (Billinghurst; Clark; Lee, 2015), while consumer experiments like Google Glass (2013) explored heads-up interfaces before being discontinued in 2023 (Ruth, 2024).

Smartphones played a critical role in scaling AR adoption. Apps like Pokémon GO (2016) introduced mainstream users to AR through camera overlays. ARKit (Apple) and ARCore (Google), launched in 2017, enabled mobile AR with motion and depth tracking (Vieyra; Vieyra, 2018).

Most recently, the line between AR and VR is blurring. Apple’s Vision Pro (announced 2023) merges high-resolution VR with passthrough AR, positioning itself as a “spatial computer.” With features like dual 4K displays and hand/eye tracking, it may represent a watershed moment for XR despite its premium price (Ruth, 2024).

As of 2025, the hardware ecosystem spans from mobile-based AR apps to advanced mixed reality headsets, forming a robust toolbox for immersive educational simulations.

2.1.2 Software Ecosystems and Frameworks:

Alongside hardware, a mature software ecosystem has enabled rapid development of immersive simulations. Modern game engines such as Unity and Unreal Engine have become the de facto platforms for AR/VR content creation. These engines provide high-performance 3D graphics rendering, physics simulation, and cross-platform deployment, greatly simplifying the creation of interactive virtual environments. Unity, for example, offers an entire XR development toolkit (with support for VR headsets and AR through packages like AR Foundation) that abstracts away device-specific details and allows de-

velopers to build an application once and deploy across multiple headsets (Atta *et al.*, 2022). Unreal Engine likewise includes integrated support for VR rendering and AR (via ARKit/ARCore plugins), making high-fidelity visualization accessible to developers in academia and industry.

For mobile AR, platform-specific frameworks are key. Apple’s ARKit (introduced in iOS 11, 2017) and Google’s ARCore (for Android, 2017) brought advanced AR capabilities to hundreds of millions of smartphones (Vieyra; Vieyra, 2018). These software development kits handle real-time tracking of the device’s position, surface detection, lighting estimation, and more, allowing apps to place and persist virtual objects in the user’s environment. Thanks to ARKit/ARCore, an educator can deploy an AR simulation on standard tablets or phones – for instance, letting students point an iPad at a textbook and see 3D molecules or physical field lines appear “attached” to the pages. On the web, the WebXR API has emerged as a W3C standard enabling AR and VR experiences to run directly in web browsers using JavaScript (World Wide Web Consortium, 2021). WebXR (successor to earlier WebVR/WebAR efforts) allows an immersive educational module to be accessed with a simple URL, lowering the barrier to entry (no app install required) and ensuring compatibility across different devices (from VR headsets to phones). This is particularly relevant for broad educational deployments, where web-based delivery can be more practical. Complementing these are various supporting frameworks: for example, libraries for spatial mapping, hand tracking, and user interaction (e.g. Microsoft’s Mixed Reality Toolkit for Unity, or Vuforia for image-target AR) which provide higher-level tools for common AR/VR interactions. There are also open standards like OpenXR (released by the Khronos Group in 2019) that unify the interface to VR/AR hardware – a developer can write code once against OpenXR and run on any compliant headset (Oculus, SteamVR, Windows Mixed Reality, etc.), which is increasingly adopted by engines and platforms. In summary, the software landscape – from powerful 3D engines to AR phone toolkits and web standards – has matured to a point that immersive educational simulations can be built with relatively modest effort compared to a decade ago. This thesis will leverage these tools to construct its simulation system, ensuring it is built on proven, widely supported technology.

2.1.3 Use Cases in Education:

AR and VR have shown strong potential to enhance learning, particularly in subjects involving abstract or spatial concepts. Their core strength lies in making the invisible visible and the abstract tangible. In physics education, for instance, VR has helped students visualize and manipulate 3D vectors, improving understanding of vector addition and spatial relationships (Campos; Hidrogo; Zavala, 2022). Studies show that such immersive

tools can boost engagement and deepen comprehension of abstract STEM topics like electromagnetism or geometry through interactive, risk-free exploration (Campos; Hidrogo; Zavala, 2022; Johnson-Glenberg, 2018).

In astronomy and aerospace, where scales are far beyond human experience, immersive technologies offer unique advantages. VR enables virtual field trips through space — letting students stand on Mars or orbit planets — providing an intuitive grasp of scale and distance. Learners can explore the solar system with accurate proportions, making complex spatial relationships (like planetary distances or ring sizes) more comprehensible (Atta *et al.*, 2022). Astrophysical phenomena such as orbital mechanics and black hole dynamics are also made more accessible through interactive VR visualizations.

In aerospace engineering, VR and AR are increasingly used for hands-on training. Beyond traditional flight simulators, modern VR platforms allow students to perform simulated pre-flight inspections, engine maintenance, or spacecraft docking. Vaughn College, for example, uses VR for aviation trainees to practice inspecting and assembling parts, reinforcing mechanical familiarity before real-world exposure. Similarly, Atta *et al.* (2022) created a virtual “space lab” where students assemble a CubeSat in a simulated cleanroom, boosting their understanding of subsystem configuration through direct interaction and gamified tasks (Atta *et al.*, 2022).

AR complements this by overlaying digital instructions on real-world hardware. NASA’s Project Sidekick exemplifies this: astronauts use HoloLens headsets aboard the ISS to receive real-time, spatially anchored maintenance guidance (NASA, 2015). In classrooms, AR enables students to interact with 3D models of rockets or overlay CAD designs onto physical parts, enriching theoretical lessons with live, contextual visualization (Atta *et al.*, 2022; Milgram; Kishino, 1994).

2.1.4 Embodiment, Interaction, and Spatial Cognition:

A recurring theme in the educational use of AR/VR is the role of embodied and spatial learning. Immersive technologies engage the human sensorimotor system – users move their bodies to navigate virtual spaces, use gestures to interact with virtual objects, and perceive environments at true scale. This physicality supports cognitive processing by leveraging innate spatial reasoning and muscle memory. The theory of embodied cognition holds that learning is grounded in the body’s interactions with its environment, and AR/VR extend this principle digitally. Johnson-Glenberg (2018) highlights the pedagogical value of 3D gestures: when learners rotate a virtual object or walk through a graph, they build stronger memory links (Johnson-Glenberg, 2018). Her research shows that full-motion VR, where body movements align with abstract concepts, can deepen understanding and recall. Complementary studies (e.g., Liu *et al.*, 2020) found improved reten-

tion when students enacted phenomena physically, and also noted increased presence and agency—factors tied to motivation (Campos; Hidrogo; Zavala, 2022; Johnson-Glenberg, 2018).

Spatial cognition benefits are also well-documented. VR’s stereoscopic depth and six degrees of freedom help learners perceive complex spatial relationships, vital in subjects like anatomy, geography, and engineering. Students exploring a molecule or a solar system in VR can shift perspective freely, activating spatial memory and supporting what researchers call “situated learning” – knowledge acquired in rich spatial contexts becomes more intuitive and transferable. Campos et al. (2022), for example, found that immersive 3D interaction notably enhanced vector learning tasks requiring spatial reasoning (Campos; Hidrogo; Zavala, 2022). Similarly, in astronomy, VR’s ability to scale from the Milky Way to Earth provides concrete visualizations of abstract systems (Kersting et al., 2024).

While AR/VR offer compelling tools, they are not magic bullets – user comfort, software complexity, and thoughtful pedagogical integration remain critical (Johnson-Glenberg, 2018). Still, evidence shows that immersive simulations can enhance traditional teaching, especially for learning goals involving visualization, experimentation, or embodied experience. In the context of this thesis, the implications are clear: AR and VR form a foundational layer. They enable students to interact with simulations of aerospace systems—such as satellites or orbital dynamics—in an intuitive and experiential manner. As hardware becomes lighter and more capable, and software ecosystems more robust, immersive tools are becoming increasingly viable in education. With spatial computing platforms entering mainstream use (Ruth, 2024), AR and VR are poised not just as delivery platforms but as new paradigms for engaging with knowledge.

2.1.4.0.1 Application to This Work This thesis leverages the pedagogical strengths of immersive spatial learning through Virtual Reality as the primary modality. While the platform architecture is built on mixed reality hardware (Meta Quest 3) that supports both VR and AR passthrough modes, the educational experience emphasizes full VR immersion for the reasons detailed in Section 3.3. The literature reviewed here establishes the broader context of immersive educational technologies, while the implementation focuses specifically on VR’s capacity to place learners inside coherent spatial environments optimized for understanding orbital mechanics.

2.2 Generative Agents

Traditional software and simulations have been predominantly *deterministic*—given the same inputs, they yield the same outputs. Modern AI systems built on *genera-*

tive models, by contrast, introduce stochasticity and creativity. Large Language Models (LLMs) do not follow hard-coded rules; instead, they sample from probability distributions learned from vast textual corpora. Consequently, an LLM can produce context-dependent, varied responses rather than a single predetermined answer. This marks a paradigm shift from scripted to emergent behaviour. In recent work, advanced LLMs such as GPT-4 have even outperformed traditional reinforcement-learning agents in complex environments by reasoning through text rather than executing pre-programmed control policies (Carrasco; Rodriguez-Fernandez; Linares, 2025). While stochastic generation entails some unpredictability, it is precisely this creativity that lets *generative agents* adapt to scenarios beyond their designers' foresight.

At a conceptual level an LLM is a statistical language engine: given a textual history, it predicts the most plausible continuation one word at a time. Because it is trained on heterogeneous data, a single model can answer coding questions, analyse legal texts, or reason about orbital mechanics when prompted appropriately. This broad, generative capability underpins the rise of *LLM-powered agents* (Anthropic, 2024).

LLM-based agents are autonomous software entities that embed an LLM as their core “brain.” An agent senses its environment, reasons about goals, and acts—iteratively—until a task is complete. Industry definitions describe such an agent as “a system that uses an LLM to reason through a problem, create a plan, and execute that plan with tools” (Chen, 2023; Huang; Grady, *et al.*, 2024). The LLM supplies the reasoning; auxiliary modules provide planning, memory, and tool use (Anthropic, 2024). Crucially, the agent—not the user—controls the loop: it may decide which function to call, when to revise a plan, or whether to request clarification (OpenAI, 2023). Hence an agent is more than a single LLM invocation; it is a continual perceive–think–act cycle.

Architectural Components

Generative-agent designs typically comprise five interacting elements (Anthropic, 2024; Huang; Grady, *et al.*, 2024):

- **Planning and reasoning.** The agent decomposes high-level goals into actionable steps, often prompting the LLM to produce an internal plan or “chain of thought.”
- **Memory.** Short-term context (recent turns) and long-term knowledge (summaries or retrieved documents) are stored externally—e.g. in a vector database—and injected into prompts as needed.
- **Tool use and APIs.** Through structured outputs (JSON function calls, shell commands, *etc.*) the agent invokes external tools to compute, query, or effect changes in its environment (OpenAI, 2023).

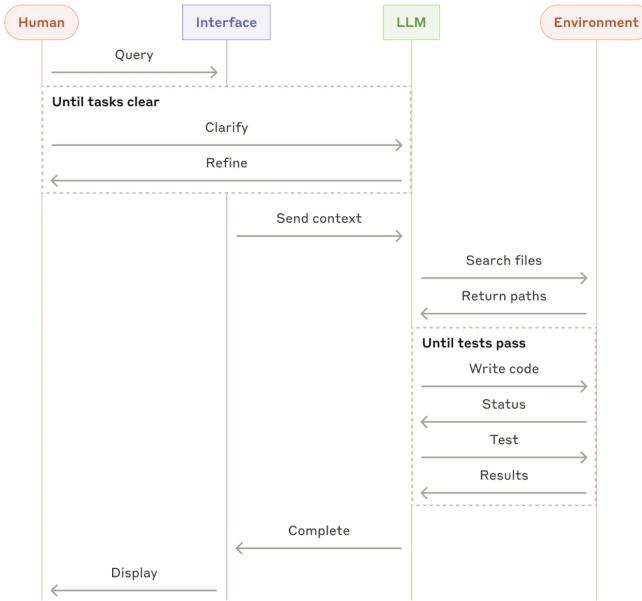


FIGURE 2.1 – End-to-end agentic workflow from Anthropic’s “Building Effective Agents.” The human issues a query through an *interface*; the LLM asks clarifying questions until the task is precise, receives contextual files, iteratively writes and tests code against the environment, and finally returns results for display (Anthropic, 2024).

- **Iterative control loop.** The agent cycles through *observe* → *reason* → *act* → *observe*, optionally reflecting or self-critiquing between steps to improve reliability.
- **Autonomy and adaptation.** Equipped with the above, the agent can switch strategies, recover from errors, and pursue its objective with minimal human micromanagement.

Applications and Relevance

- **Simulations and interactive worlds.** Park *et al.* created “Generative Agents” that populate a sandbox town with virtual characters who plan, remember, and socially interact—producing emergent storylines never scripted by the developers (Park *et al.*, 2023).
- **Aerospace guidance and control.** Carrasco *et al.* demonstrated an LLM agent piloting a spacecraft in the *Kerbal Space Program* simulation by iteratively reading textual telemetry and issuing control actions, matching classical controllers without explicit orbital equations (Carrasco; Rodriguez-Fernandez; Linares, 2025).
- **Legal reasoning.** Harvey AI equips law-firm associates with an agent that drafts memos, retrieves precedents, and iteratively refines analyses through dialogue—illustrating agentic workflows in language-dense tasks (Chen, 2023).

- **Education.** Khan Academy's *Khanmigo* employs GPT-4 as a Socratic tutor that adapts explanations to each learner, providing hints rather than answers and thereby personalising study sessions at scale (Academy, 2023).

2.3 Orbital Mechanics: The Physics of Celestial Motion

The intuitive, visual understanding of orbital motion is a primary objective of this project. While the generative agent handles the underlying calculations, a firm grasp of the governing principles is essential to frame the simulation's logic and appreciate its educational value. Orbital mechanics is the study of the motion of bodies under the influence of gravity. For missions in Earth's orbit and for interplanetary trajectories, the foundational principles discovered by Isaac Newton and Johannes Kepler provide a remarkably accurate framework for describing and predicting these celestial paths. This section outlines the core concepts that form the physical and mathematical basis of the simulation system, focusing on the specific orbital regimes and parameters that the platform enables users to explore: circular orbits, elliptical orbits, and escape trajectories.

2.3.1 Newtonian Gravity and the Two-Body Problem

At the heart of all orbital motion lies gravity. In the 17th century, Sir Isaac Newton formulated the Law of Universal Gravitation, stating that any two bodies attract each other with a force proportional to the product of their masses and inversely proportional to the square of the distance between them (Curtis, 2020). This is expressed mathematically as:

$$F = G \frac{m_1 m_2}{r^2}$$

where F is the gravitational force, G is the gravitational constant ($6.674 \times 10^{-11} \text{ N}\cdot\text{m}^2/\text{kg}^2$), m_1 and m_2 are the masses of the two bodies, and r is the distance between their centers.

When applied to a satellite orbiting a celestial body like Earth, this law forms the foundation of the **two-body problem**. This model makes a critical simplifying assumption: it considers only the gravitational force between the satellite and the primary body (e.g., Earth), ignoring perturbations such as atmospheric drag, solar radiation pressure, and gravitational influences from other bodies like the Moon or the Sun (Vallado, 2013). While these forces are significant for high-precision, long-term trajectory prediction, the two-body model provides an elegant and highly accurate approximation for foundational analysis and educational purposes. The resulting equation of motion is:

$$\ddot{\vec{r}} + \frac{\mu}{r^3} \vec{r} = 0$$

Here, \vec{r} is the position vector of the satellite relative to the primary body, $\ddot{\vec{r}}$ is its acceleration, and μ (mu) is the standard gravitational parameter of the system. For Earth-orbiting satellites, $\mu = GM_{\text{Earth}} \approx 398,600 \text{ km}^3/\text{s}^2$, where M_{Earth} is Earth's mass.

The solution to this differential equation reveals a profound geometric truth: under the inverse-square law of gravity, the satellite's path must be a **conic section**—a circle, ellipse, parabola, or hyperbola (Curtis, 2020). Which conic section results depends on the satellite's energy and angular momentum. This elegant mathematical result means that all orbital trajectories, from the circular path of the ISS to the hyperbolic escape of Voyager, are governed by the same fundamental physics expressed through different geometric shapes.

2.3.2 Kepler's Laws and Orbital Geometry

Johannes Kepler, working in the early 17th century with observational data from Tycho Brahe, empirically discovered three laws of planetary motion that would later be shown to be direct consequences of Newtonian gravity. These laws provide the geometric and temporal framework for understanding orbits (Bate; Mueller; White, 1971):

1. **First Law (Law of Orbits):** The orbit of a planet (or satellite) around the Sun (or Earth) is an ellipse, with the central body at one focus. A circle is the special case of an ellipse where both foci coincide.
2. **Second Law (Law of Areas):** A line connecting the satellite to the central body sweeps out equal areas in equal times. This means the satellite moves faster when closer to the central body (at periapsis) and slower when farther away (at apoapsis).
3. **Third Law (Law of Periods):** The square of the orbital period is proportional to the cube of the semi-major axis. Mathematically: $T^2 \propto a^3$, or more precisely, $T^2 = \frac{4\pi^2}{\mu} a^3$. This law directly relates orbital size to orbital period, explaining why the ISS at 420 km altitude completes an orbit in 92.8 minutes while the Hubble Space Telescope at 540 km takes slightly longer at approximately 95 minutes.

Kepler's laws were empirical observations that Newton later proved mathematically from first principles. Together, they provide both the geometric intuition (ellipses, not circles, are the general case) and quantitative relationships (period depends on altitude) that govern orbital motion.

2.3.3 Orbital Regimes: Circular, Elliptical, and Hyperbolic Trajectories

The shape of an orbit is determined by the satellite's total mechanical energy—the sum of its kinetic energy (from motion) and gravitational potential energy (from position in the gravity field). This energy dictates which conic section describes the trajectory. The platform's simulation implements three fundamental orbital regimes, each representing a different energy state and mission application.

2.3.3.1 Circular Orbits: Stable Operational Platforms

A **circular orbit** occurs when the satellite's velocity is precisely calibrated so that the centripetal acceleration required for circular motion exactly matches the gravitational acceleration at that altitude. This is the special case where eccentricity $e = 0$.

For a circular orbit at radius r from Earth's center (altitude $h = r - R_{\text{Earth}}$), the required orbital velocity is given by:

$$v_{\text{circular}} = \sqrt{\frac{\mu}{r}}$$

This relationship shows that orbital speed decreases with altitude: satellites in low Earth orbit (LEO) travel faster than those in higher orbits. For example, the ISS at 420 km altitude orbits at approximately 7.66 km/s, while the Hubble Space Telescope at 540 km altitude travels at approximately 7.59 km/s—slightly slower due to its higher altitude.

Circular orbits are preferred for operational missions requiring predictable, repeating ground tracks and stable altitude. The International Space Station (420 km, 51.6° inclination) and Hubble Space Telescope (540 km, 28.5° inclination) both use circular orbits because their missions benefit from the stability and predictability of constant altitude and speed.

Kepler's Third Law directly determines the orbital period for circular orbits:

$$T = 2\pi \sqrt{\frac{r^3}{\mu}}$$

This equation explains the relationship between altitude and period. The ISS at 420 km completes 15.5 orbits per day, providing frequent revisit times for Earth observation and crew operations. Hubble at 540 km has a slightly longer period, chosen to balance orbital stability with minimizing atmospheric drag while providing optimal viewing conditions for astronomical observations.

2.3.3.2 Elliptical Orbits: Variable Altitude Trajectories

An **elliptical orbit** occurs when $0 < e < 1$, where e is the eccentricity. The satellite's altitude and speed vary continuously as it moves around the ellipse. The closest point to Earth is called **periapsis** (or perigee for Earth orbits), and the farthest point is **apoapsis** (or apogee). The size of the ellipse is characterized by the **semi-major axis** a , which is half the longest diameter of the ellipse.

The relationship between the semi-major axis, periapsis radius r_p , and apoapsis radius r_a is:

$$a = \frac{r_p + r_a}{2}$$

The eccentricity quantifies how elongated the ellipse is:

$$e = \frac{r_a - r_p}{r_a + r_p}$$

When $e = 0$, the ellipse becomes a circle ($r_p = r_a$). As e approaches 1, the ellipse becomes increasingly elongated.

The satellite's speed at any point in an elliptical orbit is given by the **vis-viva equation**, one of the most fundamental relationships in orbital mechanics (Curtis, 2020):

$$v = \sqrt{\mu \left(\frac{2}{r} - \frac{1}{a} \right)}$$

This equation reveals that orbital speed depends on both the current position r and the orbit's overall size a . At periapsis, where r is smallest, the satellite moves fastest. At apoapsis, where r is largest, it moves slowest. This speed variation is a direct consequence of Kepler's Second Law: the satellite must move faster when closer to Earth to sweep equal areas in equal times.

Elliptical orbits have important applications. Highly Elliptical Orbits (HEO) are used for communications satellites serving high-latitude regions, as the satellite spends most of its time near apoapsis with excellent visibility over polar regions. Transfer orbits between circular orbits are also elliptical, with the initial circular orbit at periapsis and the target circular orbit at apoapsis.

2.3.3.3 Hyperbolic Trajectories: Escaping Earth's Gravity

A **hyperbolic trajectory** occurs when $e \geq 1$. Unlike elliptical orbits, which are closed and periodic, hyperbolic trajectories are open curves—the spacecraft approaches Earth,

swings around it, and departs, never to return. This regime represents escape from Earth's gravitational influence.

The minimum speed required to achieve escape from Earth's surface is the **escape velocity**:

$$v_{\text{escape}} = \sqrt{\frac{2\mu}{r}}$$

At Earth's surface ($r = R_{\text{Earth}} = 6371$ km), this yields approximately 11.2 km/s. Notice that escape velocity is exactly $\sqrt{2}$ times the circular orbital velocity at the same radius—this factor of $\sqrt{2}$ represents the energy difference between a bound circular orbit and an unbound escape trajectory.

For a hyperbolic trajectory, the spacecraft's velocity at any distance is given by a modified vis-viva equation:

$$v = \sqrt{\mu \left(\frac{2}{r} + \frac{1}{a} \right)}$$

Note the sign change: for hyperbolic orbits, the semi-major axis a is defined as negative, reflecting the fact that the orbit is unbound with positive total energy.

The Voyager spacecraft exemplify hyperbolic escape trajectories. After launch and acceleration to sufficient velocity, they followed hyperbolic paths that carried them beyond Earth's sphere of influence and into interplanetary space. The platform uses Voyager's trajectory to demonstrate the transition from bound elliptical motion to unbound hyperbolic escape, illustrating the fundamental energy threshold that separates orbiting from departing.

2.3.4 Orbital Orientation: Inclination and Coverage

While the size and shape of an orbit (determined by a and e) govern its energy and geometry, the **inclination** determines the orbit's orientation in three-dimensional space. Inclination i is the angle between the orbital plane and a reference plane, typically Earth's equatorial plane. An inclination of 0 defines an equatorial orbit, while 90 defines a polar orbit that passes directly over both poles.

Inclination is not arbitrary—it is fundamentally constrained by the launch site's latitude and the physics of rotation. When a rocket launches eastward (prograde), it benefits from Earth's rotational velocity, which is maximum at the equator (~465 m/s) and decreases toward the poles. The minimum achievable inclination from a launch site is approximately equal to the site's latitude. For example, launches from Kennedy Space Center (28.5°N) can achieve inclinations of 28.5° or greater, but reaching lower inclinations would require the rocket to perform an energetically expensive plane change maneuver.

This launch constraint explains many mission orbital parameters:

- **ISS (51.6° inclination):** Designed to be accessible from both Kennedy Space Center and the Baikonur Cosmodrome in Kazakhstan (45.6°N). The 51.6° inclination allows Russian Soyuz launches from Baikonur while remaining within reasonable energy budgets for US launches.
- **Hubble (28.5° inclination):** Launched from Kennedy Space Center at the minimum possible inclination, maximizing the rotational velocity assist and minimizing fuel requirements. This low inclination also provides good sky coverage for astronomical observations while avoiding prolonged periods in Earth’s shadow.

Inclination also determines ground track coverage. An equatorial orbit ($i = 0$) never passes over polar regions. A polar orbit ($i = 90$) eventually covers the entire surface as Earth rotates beneath it. Intermediate inclinations provide a balance between coverage and launch efficiency. The platform’s mission-specific implementations demonstrate how operational requirements (crew access for ISS, astronomical visibility for Hubble) drive inclination choices.

2.3.5 The Educational Foundation for Interactive Exploration

The physics and mathematics outlined in this section—Newtonian gravity, Kepler’s laws, the vis-viva equation, and the geometric properties of conic sections—form the computational foundation of the simulation platform. More importantly, they represent the conceptual framework that users explore through embodied interaction in virtual reality.

Traditional orbital mechanics education presents these concepts through equations on paper and two-dimensional diagrams. Students memorize formulas and solve problems numerically, but the intuitive, spatial understanding of why the ISS orbits at 7.66 km/s or why the Hubble telescope requires a specific altitude and inclination often remains elusive. The three-dimensional geometry of an inclined orbit, the speed variation along an ellipse, and the meaning of escape velocity are fundamentally spatial phenomena that are difficult to internalize from textbooks alone.

The platform’s approach inverts this pedagogy. Users begin not with equations but with questions and curiosity: "Show me the ISS orbit." "Why does Hubble orbit where it does?" "How did Voyager leave Earth?" The generative agent translates these natural language queries into the precise orbital parameters described in this section—altitude, eccentricity, inclination—and the Unity simulation engine renders the resulting trajectories as visible, three-dimensional curves in space. Users inhabit the orbital environment,

observing how the ISS's 420 km circular orbit compares to Hubble's 540 km orbit, seeing the ellipse stretch as eccentricity increases, watching the hyperbolic escape path diverge from closed elliptical motion.

This section has established the theoretical foundation that makes such exploration both accurate and meaningful. The circular orbits users create are governed by $v = \sqrt{\mu/r}$. The elliptical orbits follow the vis-viva equation. The hyperbolic escapes exceed $v_{\text{escape}} = \sqrt{2\mu/r}$. The platform's educational value rests on this foundation: it translates rigorous astrodynamics into intuitive visual experience, enabling users to build genuine understanding of orbital mechanics through guided exploration rather than rote memorization.

3 Methodology

3.1 Design Philosophy and Approach

The development of this project is fundamentally an exploratory research endeavour into a new paradigm of human-computer interaction for educational purposes. Given the innovative and complex nature of integrating generative AI, immersive mixed reality, and embodied interfaces, a rigid, waterfall-style development plan would be inappropriate. Instead, the methodology is guided by a philosophy that embraces iteration and modularity to navigate the technical challenges and discovery process inherent in such work.

The approach is defined by three core principles:

- **Prototype-Driven:** The primary goal is the creation of a functional prototype that demonstrates the feasibility and potential of the proposed system. This approach prioritizes implementing the core functionalities of the user experience over exhaustive feature development, allowing for tangible and testable results that can validate the project's central thesis.
- **Iterative Development:** The project will be built in iterative cycles, following a process of building a core feature, testing its performance and usability, and refining it based on the results. This allows for flexibility in the implementation details, acknowledging that the optimal solutions for agent prompting and user interaction will be discovered and improved upon throughout the development lifecycle.
- **Modular Architecture:** The system is designed as a collection of distinct yet interconnected modules: the generative agent (the "brain") and the simulation and visualisation engine (the "world"). This modularity, a key objective of this project, makes the complex system manageable, facilitates parallel development and testing of components, and ensures the final architecture is extensible for future work.

These principles guided a four-phase development strategy. Phase 1 established the conversational agent in the Hub environment with circular and elliptical orbit creation

capabilities. Phase 2 implemented three Mission Spaces (ISS, Hubble, Voyager) with specialist agents and scene transition tools. Phase 3 integrated bidirectional voice through ElevenLabs (Scribe v2 for speech-to-text, TTS for character synthesis). Phase 4 deployed the complete system to Meta Quest 3 with immersive VR visualization. Each phase produced a testable, working system that integrated seamlessly with subsequent development without requiring architectural changes.

3.2 System Architecture and Data Flow

The platform architecture comprises two integrated spaces: the "Hub" (Mission Control) where users create custom orbits through conversation, and "Mission Spaces" (ISS, Hubble, Voyager) where specialists demonstrate real missions. Users navigate between these environments via voice commands (`route_to_mission`, `return_to_hub`), with conversational context preserved across transitions.

The system separates conversational intelligence (OpenAI GPT-4.1) from spatial visualization (Unity 3D), connected through a tool-calling interface. The agent interprets natural language, selects appropriate tools (`create_circular_orbit`, `create_elliptical_orbit`, `set_simulation_speed`, `pause_simulation`, `reset_simulation_time`, `clear_orbit`, `route_to_mission`, `return_to_hub`), and Unity executes the corresponding physics calculations and rendering.

3.2.1 Interaction Flow

Each user interaction follows a nine-step cycle from speech input to audio output:

1. User speaks request via Quest 3 microphone
2. ElevenLabs Scribe v2 transcribes audio to text
3. OpenAI GPT-4.1 interprets intent and selects appropriate tool
4. Unity's `ToolExecutor` validates parameters and invokes corresponding C# method
5. `SceneTransitionManager` loads Mission Space (for `route_to_mission`); `OrbitController` calculates trajectory using vis-viva equation and renders visualization (for orbit creation)
6. Tool execution result returns to LLM as feedback
7. Agent generates educational response based on tool result and conversation context

8. ElevenLabs TTS synthesizes character-specific voice (Mission Control's authority, ISS specialist's precision, Voyager specialist's Sagan-like wonder)
9. Audio plays through Quest 3 spatial audio system

Conversation history is maintained across turns and scene transitions, enabling contextual dialogue. Users control all aspects—orbit creation, mission navigation, simulation parameters—through voice alone.

3.3 Mixed Reality Design Rationale

While the Quest 3 supports both VR and AR passthrough, the platform emphasizes immersive VR for pedagogical reasons. Orbital mechanics involves scales (420 km for ISS, 35,786 km for geostationary orbit) incompatible with domestic spaces—AR overlays would show trajectories passing through walls and furniture, creating perceptual friction between room-scale and cosmic-scale contexts. VR isolation places users within the orbital environment itself, establishing coherent spatial context where Earth floats in space and trajectories exist in their natural domain. The architecture supports AR passthrough for collaborative learning, museum installations, or classroom demonstrations where physical anchoring adds value, but the primary experience prioritizes the modality best suited to the content.

3.4 Technical Implementation

The system implements conversational AI (OpenAI GPT-4.1 for reasoning, ElevenLabs for voice synthesis/transcription) interfaced with Unity 3D physics simulation. The agent embodies four characters—Mission Control in the Hub, plus three mission specialists (ISS crew perspective for LEO operations, Hubble engineer for telescope mission design, Voyager/Sagan persona for interplanetary trajectories)—each with distinct voice profiles and expertise areas. Unity implements two-body orbital physics in C#, calculates trajectories via vis-viva equation, and renders visualizations on Quest 3. The platform supports circular orbit creation (160-35,786 km altitude) and elliptical orbit creation (periapsis/apoapsis 160-100,000 km), with inclination constrained to 0-180°.

3.5 Development and Version Control

The project follows systematic development practices using GitHub for version control. All source code—including Unity C# scripts, prompt templates, and configuration

files—is tracked in a central repository, providing complete history of changes and enabling experimental work through branching without compromising the main project stability. This systematic approach aligns with the iterative development philosophy, where each development cycle’s progress is documented and preserved.

3.6 Validation and Dissemination Strategy

The platform’s educational effectiveness and technical implementation are validated through demonstration of complete interaction scenarios across all mission spaces. The validation approach prioritizes real-world usage patterns: users exploring orbital mechanics through natural conversation, navigating between Hub and Mission Spaces, and experiencing the full voice-driven VR workflow. The complete implementation is released open-source on GitHub, enabling community validation, adaptation, and extension. Users provide their own API keys for OpenAI and ElevenLabs services, ensuring accessibility without imposed service costs. Comprehensive documentation covers Unity configuration, Quest 3 deployment, and integration procedures.

3.7 Core Module Implementation

This section describes how the platform’s four primary modules—conversational agent, orbital physics simulation, voice integration, and virtual reality environment—work together to create an immersive educational experience for learning orbital mechanics. Each subsection explains the educational rationale behind key design decisions and how they support the learning objectives established in Section 3. Complete technical specifications are provided in Appendices A–D.

3.7.1 Agent System Implementation

The conversational agent system removes the traditional barrier between learning intent and technical execution by enabling natural language control of orbital simulations. Learners can express goals like “Create an orbit matching the ISS” or “Show me a highly elliptical orbit” without needing to understand programming, coordinate systems, or simulation APIs. This design decision directly addresses a core challenge in physics education: allowing students to focus on conceptual understanding rather than technical syntax.

Following the tool-calling architecture outlined in Section 2.2, the system implements GPT-4.1 with a structured tool-calling framework that interprets user requests and invokes validated simulation commands. When a learner asks to create an orbit, the agent trans-

lates natural language into precise physics parameters (altitude, inclination, eccentricity), executes the orbital calculation, and explains the result in educational terms. Critically, the agent disambiguates between orbital velocity (the speed required to maintain a specific orbit, calculated from physics) and simulation time speed (how fast the visualization plays back)—a common source of confusion that this explicit separation prevents.

The platform embodies two agent archetypes: Mission Control (at the Hub) focuses on orbit creation and simulation control, while three mission specialists (ISS, Hubble, Voyager) provide mission-specific educational context when learners navigate to dedicated Mission Spaces. This dual-character design supports two learning modes: hands-on experimentation at the Hub, and contextual deepening through mission-specific dialogue. Conversation history persists across scene transitions, enabling learners to ask follow-up questions like “What was the altitude of the orbit I just created?” after switching contexts—supporting iterative, exploratory learning patterns.

Technical implementation details, including prompt architecture, tool schemas, context management algorithms, and API integration specifications, are documented in Appendix A.

3.7.2 Orbital Physics Simulation

The orbital physics engine translates altitude specifications into velocity requirements automatically, making visible a fundamental relationship that students often struggle to grasp: that orbital speed is not arbitrary but determined by altitude through gravitational physics. When a learner requests “an orbit at 420 km like the ISS,” the system calculates the required velocity (7.66 km/s) using the vis-viva equation and displays both values together. This automatic calculation prevents a common misconception—that higher orbits move faster—by immediately showing that geostationary satellites at 35,786 km altitude actually travel slower (3.07 km/s) than low Earth orbit satellites, despite their greater distance.

The simulation implements two-body Keplerian mechanics (Section 2.3) with physically accurate trajectory calculations for circular and elliptical orbits. Visual trajectories render as continuous curves in VR space, allowing learners to observe geometric properties directly: circular orbits maintain constant radius, while elliptical orbits visually demonstrate eccentricity through their oblong shape. Scale compression maps Earth’s 6,371 km radius to a comfortable VR viewing volume while preserving proportional relationships—the ISS appears at 6.6% of Earth’s radius above the surface, matching the real ratio—enabling learners to develop accurate spatial intuition about orbital altitudes without being overwhelmed by vast scales.

Critically, all physics calculations occur in real units (km, km/s) before conversion to rendering space, ensuring that displayed values match published orbital data for ISS, Hubble, and other missions. This fidelity allows learners to verify simulation results against authoritative sources, building confidence in the educational tool. Complete physics implementation, including vis-viva equation derivations, scale compression algorithms, and trajectory visualization methods, appears in Appendix B.

3.7.3 Voice Integration Pipeline

Voice interaction addresses a practical constraint of immersive VR environments (Section 2.1): hands holding controllers cannot easily type, and virtual keyboards break immersion. The system implements bidirectional speech through push-to-talk input (Quest 3 controller A button) and synthesized character voices, enabling learners to engage in natural spoken dialogue while manipulating 3D orbital visualizations. This hands-free modality supports exploratory learning patterns where students voice hypotheses (“What happens if I increase the altitude?”), observe results, and refine understanding through iterative questioning—a cognitive process difficult to sustain when switching between physical keyboards and immersive VR.

Each agent character embodies a distinct voice: Mission Control speaks with authoritative encouragement at the Hub, while mission specialists (like Anastasia, the ISS expert) adopt personalities aligned with their educational roles—professional, technical, and approachable. This character differentiation serves pedagogical purposes beyond engagement: learners develop associative memory between voice identity and knowledge domain, reinforcing context switching as they navigate between experimental workspace (Hub) and mission-specific deepening (ISS, Hubble, Voyager spaces). Voice synthesis occurs within 1–3 seconds of agent response generation, maintaining conversational flow without perceptible delays that would disrupt the learning dialogue.

The push-to-talk mechanism balances spontaneity with intentionality: learners explicitly signal when they wish to speak, preventing accidental voice activation while preserving the natural rhythm of conversation. This design choice emerged from recognizing that educational dialogue differs from commercial voice assistants—students need time to think between questions, and the platform should not interpret silence as disengagement. Technical details of speech-to-text processing, audio synthesis parameters, and character voice management appear in Appendix C.

3.7.4 Virtual Reality Environment

As discussed in Section 2.1, the virtual reality environment transforms abstract orbital mechanics into spatial experiences that leverage human depth perception and proprioception. Orbits exist as three-dimensional curves that learners can walk around, crouch beneath, and observe from multiple vantage points—building geometric intuition impossible to achieve through 2D screens or static diagrams. Seeing an elliptical orbit’s eccentricity from different angles, or observing how inclination tilts the orbital plane relative to Earth’s equator, engages spatial reasoning faculties that support conceptual understanding of orbital geometry.

The platform deploys to Meta Quest 3, a standalone VR headset enabling tetherless movement around orbital visualizations without PC connection constraints. Maintaining 90 Hz stereoscopic rendering ensures visual comfort during extended learning sessions, preventing the nausea and fatigue that would undermine educational effectiveness. This frame rate requirement drove architectural decisions throughout the implementation: single-pass instanced rendering reduces GPU overhead, texture compression minimizes memory bandwidth, and asynchronous scene loading prevents visible stuttering during navigation between Hub and Mission Spaces.

The multi-scene architecture supports distinct learning contexts: the Hub provides an experimental workspace for orbit creation and manipulation, while three Mission Spaces (ISS, Hubble, Voyager) offer focused environments for deepening understanding of specific missions. Scene transitions preserve conversation history and simulation state, allowing learners to seamlessly shift between hands-on experimentation and contextual exploration. Spatial UI elements render in 3D world space rather than head-locked overlays, maintaining presence and spatial grounding while providing necessary information—mission elapsed time, simulation speed, and dialogue responses appear as objects in the environment rather than disconnected interface chrome.

Technical specifications for Quest 3 deployment, including Android build configuration, input system implementation, stereo rendering pipeline, and performance optimization strategies, are detailed in Appendix D.

4 Results and Demonstration

This chapter demonstrates the complete functional platform through the lens of an actual user learning journey. Rather than presenting isolated features, each section follows the narrative arc of the demonstration session, showing how a learner explores orbital mechanics concepts through curiosity-driven dialogue with AI specialists, hands-on orbit creation, and immersive VR visualization.

The demonstration validates all six specific objectives established in Section 1.3 through integrated scenarios:

- **Objective #1 (Physically accurate simulation):** Orbital velocities and periods match published mission data (ISS: 7.66 km/s at 422 km altitude)
- **Objective #2 (Natural language interpretation):** User questions like “What’s a good altitude?” trigger appropriate tool selection and specialist routing
- **Objective #3 (Multimodal VR interaction):** Voice commands, spatial presence, and Quest 3 stereoscopic rendering enable embodied learning¹
- **Objective #4 (Real-time coherence):** Scene transitions, voice synthesis, physics calculations, and rendering execute reliably without interruption
- **Objective #5 (Modular architecture):** Hub and three Mission Spaces (ISS, Hubble, Voyager) operate as distinct scenes sharing common infrastructure
- **Objective #6 (Open-source delivery):** Complete source code, documentation, and deployment guides enable community validation

The chapter structure follows the demonstration’s pedagogical pattern: *question* → *specialist guidance* → *hands-on creation* → *visual observation* → *iterative refinement*. This complete cycle repeats twice for circular orbits (Section 4.2) and elliptical orbits (Section 4.3), demonstrating experiential learning through hands-on creation. The escape trajectory section (Section 4.4) demonstrates a complementary mode: conceptual dialogue

¹This demonstration focuses on VR immersion. AR passthrough capabilities exist in the codebase (`ARHub.unity`) but are not exercised in this session.

without hands-on creation, showing the platform’s flexibility to support both experiential and theoretical learning depending on pedagogical context.

4.1 Entering the Orbital Environment

The platform’s learning journey begins when the user puts on the Quest 3 headset and enters the virtual environment. This onboarding sequence—cutscene introduction, Hub arrival, and first interaction—establishes spatial context, introduces the conversational agent, and demonstrates the voice-driven interface that will mediate all subsequent exploration.

4.1.1 Opening Sequence: Narrator Introduction

The experience opens with a cinematic sequence. The user’s viewpoint begins in deep space, far from Earth, while a narrator’s voice provides context:

NARRATOR: Above us, thousands of satellites trace perfect arcs—testaments to gravity, velocity, and the geometry of motion. You’re about to join that tradition. I’ll guide the physics; you’ll design the orbits. Together, we’ll explore how spacecraft navigate the cosmos.

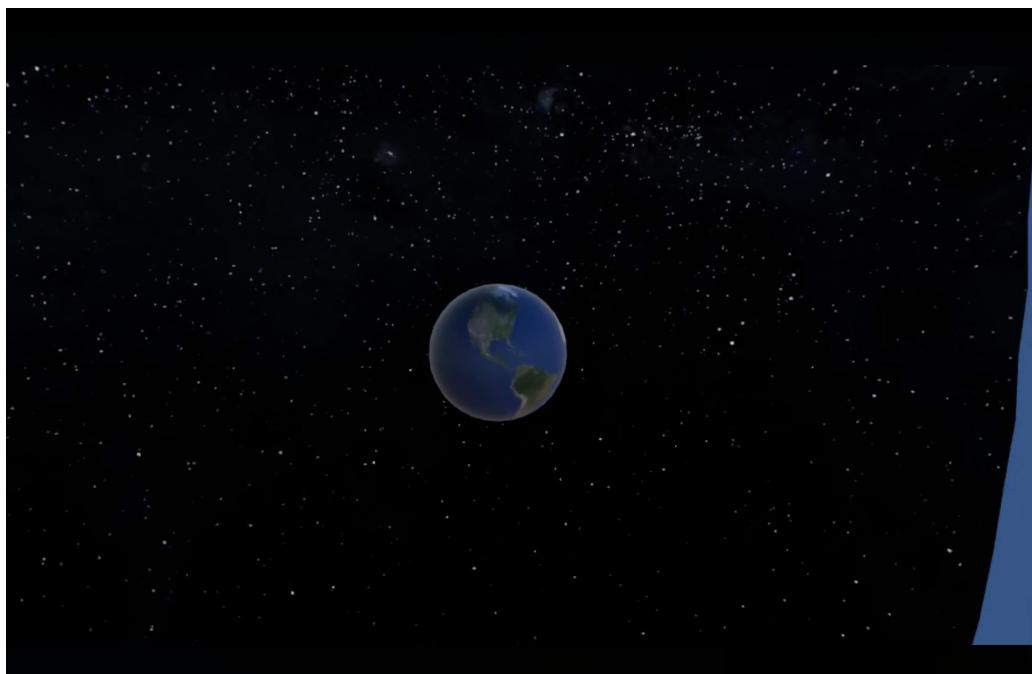


FIGURE 4.1 – Opening cutscene frame showing Earth from deep space. The camera begins at this distant perspective before zooming toward the Hub’s orbital position.

As the narrator speaks, the camera moves smoothly toward Earth over 28.7 seconds. This sequence serves two purposes: it establishes the cosmic scale of the environment the user is about to inhabit, and it frames the interaction paradigm. The phrase “I’ll guide the physics; you’ll design the orbits” positions the learner as active participant rather than passive observer—a signal that this platform prioritizes exploration over instruction.

The camera animation is managed by the `ExperienceManager`, which coordinates the zoom from deep space to Hub position using Unity’s Cinemachine interpolation. The transition occurs asynchronously while the narrator’s synthesized voice continues, ensuring smooth audio-visual integration.

4.1.2 Hub Environment Arrival: Spatial Presence and Agent Welcome

When the cutscene concludes, the user arrives at the Hub—Mission Control’s orbital workspace. The transition from cinematic observer to embodied presence is immediate: the user now floats in space, surrounded by stars, with Earth directly ahead.

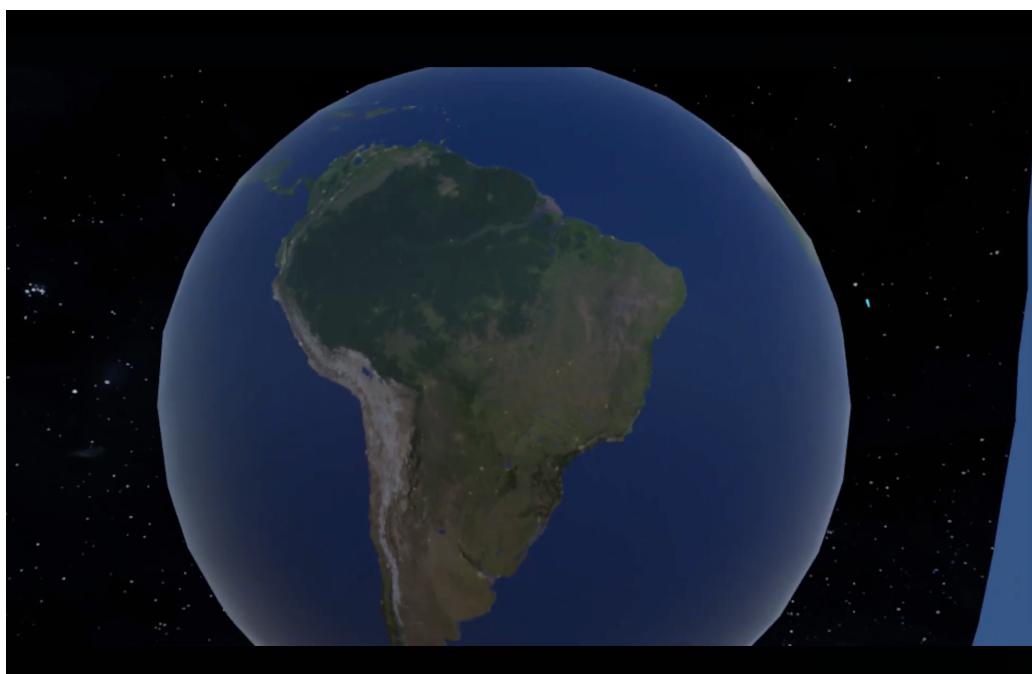


FIGURE 4.2 – Hub environment from user perspective. Earth floats in stereoscopic three-dimensional space, rendered with 8K textures to provide visual fidelity while maintaining 90 Hz frame rate on Quest 3.

CAPCOM, the Mission Control agent, greets the user with both spatial and interactional orientation:

CAPCOM: Welcome to Mission Control. You’re floating above Earth. Use your controllers to drift closer. Circle around and watch the continents slide beneath you, oceans catching light, night falling while dawn breaks on the

other side. I'm Capcom. We can design orbits together, build trajectories, see how satellites actually move—or I can route you to three specialists:

- a crew member aboard the ISS who lives this work every day,
- an engineer who built Hubble to see what we couldn't from down here,
- and a space enthusiast who dreams about what Voyager means beyond everything.

When you're ready to speak, press A on your right controller. I'll be listening. Enjoy the view—that blue marble beneath you is home.

CAPCOM's welcome accomplishes several tasks simultaneously. First, it orients the user spatially: "you're floating above Earth." Second, it demonstrates the VR locomotion—users can move within their Guardian boundary and use controllers to navigate, allowing them to observe Earth from any angle. Third, it introduces the push-to-talk mechanism: "press A on your right controller." This voice-first interface is essential in VR, where keyboards are impractical. Finally, CAPCOM previews the three Mission Spaces—ISS, Hubble, and Voyager—that represent the modular architecture enabling context-specific learning.

The voice itself conveys character. CAPCOM speaks with professional warmth, balancing authority with encouragement. This tone, shaped through ElevenLabs synthesis with a carefully selected voice profile, distinguishes Mission Control from the three specialists who will appear later with their own vocal identities and expertise.

4.1.3 Satellite Discovery: First Interaction

After orienting to the environment, the user notices a spacecraft model floating near Earth's surface. This three-dimensional satellite will serve as the visual anchor for all orbital trajectories created during the session. The user moves closer to examine it, then presses the A button to speak:

USER: Hey, Capcom. I see we have a satellite here. Is this the satellite that we can build?

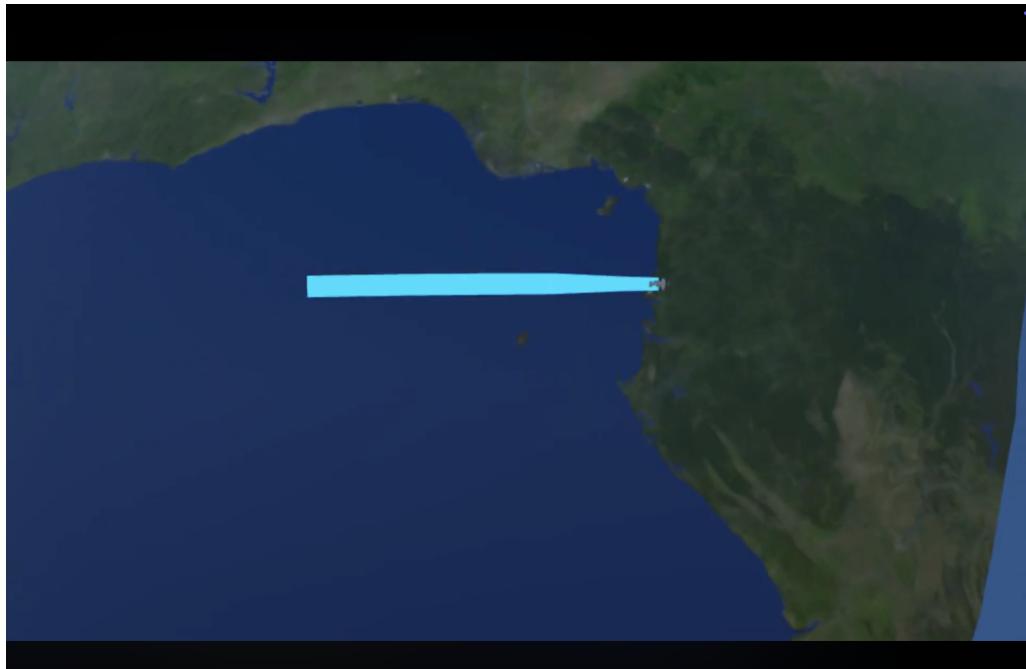


FIGURE 4.3 – Close-up view of the simulation satellite in the Hub environment. The user has navigated their perspective to examine the spacecraft model.

This simple question exercises the complete voice pipeline for the first time. When the user presses and holds the A button, the Quest 3 microphone captures audio. Upon release, the recording transmits to ElevenLabs Scribe v2 for transcription. Within 1–2 seconds, the text returns: “Hey, Capcom. I see we have a satellite here. Is this the satellite that we can build?” The transcribed text passes to OpenAI’s GPT-4.1, which generates a contextual response. That response synthesizes back into audio through ElevenLabs TTS, now in CAPCOM’s voice, and plays through the headset’s spatial audio system.

The agent responds in conversational mode—no tool execution is required, only dialogue:

CAPCOM: [Music] Hello from Mission Control. Yes—this is your simulation satellite. We can build orbits by setting circular or elliptical paths. Tell me either the altitude for circular, or periapsis and apoapsis for elliptical, and I’ll get it launched.

CAPCOM’s answer does more than confirm the satellite’s identity. It teaches vocabulary: “altitude” defines circular orbits, while “periapsis and apoapsis” specify elliptical orbits. This disambiguation will support understanding when the user begins creating trajectories in subsequent sections. The phrase “you can build” reinforces agency—the learner is positioned as designer, not spectator.

The entire interaction cycle—button press, speech, transcription, agent reasoning, synthesis, playback—completes in approximately 4–6 seconds. This latency, dominated

by external API calls to ElevenLabs and OpenAI rather than platform inefficiency, proves acceptable for educational dialogue where thoughtful responses matter more than instant reactivity.

4.1.3.0.1 Onboarding Complete This first exchange, though brief, establishes the platform’s interface paradigm. The user now understands where they are (orbital space above Earth), who they’re speaking with (CAPCOM at Mission Control), how to communicate (push-to-talk voice), what they can do (design circular and elliptical orbits, consult mission specialists), and what object they’ll be controlling (the simulation satellite). With these foundations in place, exploratory learning can begin.

Section 4.2 follows the first complete pedagogical cycle: the user asks a conceptual question about orbital altitude, consults the ISS specialist for guidance, returns to create a circular orbit, and observes its motion through VR visualization.

4.2 Learning Through Mission Specialist Dialogue: The ISS Circular Orbit

This section demonstrates the platform’s complete pedagogical cycle: a learner asks a conceptual question, the agent routes to a mission specialist for expert guidance, the user returns to create an orbit based on that knowledge, and observes the resulting physics in immersive VR. This cycle embodies the learning paradigm described in Section 1.2: spatial, conversational, and adaptive education that transforms abstract orbital mechanics into tangible experience.

The demonstration follows a natural learning progression: the user wants to create a circular orbit but doesn’t know what altitude to choose. Rather than providing a formula or arbitrary number, CAPCOM connects the user with ANASTASIA—an ISS crew member who explains real-world altitude selection through operational context. This specialist consultation transforms “altitude” from an abstract parameter into a meaningful engineering decision with practical trade-offs.

4.2.1 The Learner’s Question: Seeking Conceptual Foundation

Having learned that circular orbits require an altitude parameter (Section 4.1.3), the user recognizes a knowledge gap and asks a foundational question:

USER: Cool. Let’s start with a circular orbit. I need an altitude—but what’s a good altitude? What determines a good altitude?

This question represents genuine curiosity-driven learning. The user is not asking for procedural instruction (“how do I create an orbit?”) but conceptual understanding (“what makes an altitude good?”). This distinction is pedagogically significant: the learner seeks not just to operate the simulation but to understand orbital design principles.

CAPCOM’s agent reasoning (GPT-4.1 with `toolSelectionPrompt` from Appendix A) analyzes this input and recognizes:

- The question is *educational*, not *operational*—no tool execution needed yet
- The topic (altitude selection for circular orbits) falls within the ISS specialist’s knowledge domain
- Routing to a real-mission context will provide more meaningful learning than a generic physics explanation

The agent selects the `route_to_mission` tool with parameter `mission_id = "ISS"` and prepares a routing announcement:

*CAPCOM: [Music] Connecting you to the ISS mission for real-world examples.
Stand by.*

This response demonstrates several sophisticated capabilities:

- **Intent recognition:** Agent understood that “what determines a good altitude?” requires educational context, not simulation control
- **Specialist selection:** Correctly matched the question topic (LEO circular orbit altitude) to ISS expertise
- **Contextual routing:** The phrase “real-world examples” sets expectations that the user will receive mission-specific, applied knowledge

The tool execution triggers `ToolExecutor.ExecuteRouteToMission()` (Appendix A), which invokes `SceneTransitionManager.TransitionToMission("ISS")`. Before the scene transition begins, the routing context is preserved in `MissionContext.Instance` (Section 3.2), storing that the user arrived with the question “what determines a good altitude?”—ensuring ANASTASIA can tailor her introduction appropriately.

4.2.2 Scene Transition to ISS Mission Space: Modular Architecture in Action

The scene transition from Hub to ISS Mission Space demonstrates Objective #5 (modular architecture) and Objective #4 (real-time coherence). During the 4-second asynchronous load:

- **Visual feedback:** ISS mission logo displays as transition overlay, providing visual continuity during the load
- **Audio continuity:** Background music persists through `ExperienceManager.DontDestroyOnLoad` (Section 3.2)
- **Context preservation:** `MissionContext` maintains the conversation history (last 10 exchanges) and routing reason
- **Character preparation:** `MissionSpaceController` (Section 3.4) configures ANAS-TASIA's voice profile and specialist prompt

The user experiences no interruption in presence—VR immersion continues smoothly as the environment changes from Hub's open orbital workspace to ISS Mission Space's pre-built orbit visualization.



FIGURE 4.4 – ISS Mission Space environment after scene transition. The pre-configured ISS orbit (420 km altitude, 51.6° inclination) is visible as a cyan trajectory. ANASTASIA, the ISS specialist, greets the user with context awareness, knowing they came to learn about altitude selection. This environment demonstrates the `MissionConfig ScriptableObject` system (Appendix A) that defines mission-specific orbits, characters, and knowledge domains.

Upon scene load completion, `MissionSpaceController.Start()` executes the initialization sequence (lines 44-79 of `MissionSpaceController.cs`):

1. Creates the ISS circular orbit (420 km, 51.6° inclination) via `OrbitController.CreateCircularOrbit()`
2. Updates `MissionContext.Instance.currentLocation = "ISS"`
3. Sets ANASTASIA’s voice profile (ElevenLabs voice ID: ZF6FPAbjXT4488VcRRnw)
4. Triggers specialist introduction after 2.5-second delay (allowing for visual fade-in)

ANASTASIA’s introduction demonstrates context-aware greeting generation:

ISS_CREW (ANASTASIA): [Music] Hello from the ISS. For circular LEO orbits, engineers balance drag, launch energy, accessibility, and mission needs. The ISS flies around 420 km—high enough to limit atmospheric drag so we don’t slow down too quickly, but not so high that crew/cargo launches become inefficient. It’s also reachable from multiple launch sites. Too low: frequent reboosts. Too high: launches get harder and more expensive.

4.2.3 Specialist Educational Dialogue: Real-World Context and Engineering Trade-offs

ANASTASIA's explanation demonstrates the educational quality of specialist knowledge integration. Rather than citing the vis-viva equation ($v = \sqrt{\mu/r}$ from Section 2.3), she provides *applied engineering rationale*:

Trade-off Analysis:

- **Lower bound constraint:** “Too low: frequent reboosts” references atmospheric drag effects below 300 km
- **Upper bound constraint:** “Too high: launches get harder and more expensive” reflects launch vehicle Δv budget limitations
- **Operational optimum:** “420 km” emerges as a balanced compromise, not an arbitrary choice
- **Accessibility consideration:** “reachable from multiple launch sites” explains why ISS uses 51.6° inclination (Section 2.3—accommodates Baikonur launches)

This explanation transforms altitude from a numerical input into an *engineering decision* with multidimensional constraints. The learner now understands that “good altitude” depends on mission priorities—crew safety, launch accessibility, observational requirements, and operational costs.

The user continues the conversation with follow-up questions, demonstrating natural dialogue flow:

USER: Would it be absurd to choose the same altitude as the ISS?

ISS_CREW (ANASTASIA): Not absurd at all—quite sensible. 420 km is a proven compromise. Many Earth-observation and crewed missions use similar altitudes for the same reasons. You'd be following decades of flight heritage.

This exchange validates the learner's reasoning and reinforces that ISS's altitude is not unique but representative of a broader LEO operational regime. ANASTASIA's response uses encouraging language (“quite sensible,” “flight heritage”) that builds confidence while educating.

USER: Is the ISS circular?

ISS_CREW (ANASTASIA): Nearly. The path is close to circular with small variations (drag, gravity field). We keep it tight with periodic reboosts. Near-circular simplifies rendezvous and ops.

This technical clarification introduces realistic complexity: the ISS’s orbit is *nearly* circular, not perfectly so, due to perturbations. The mention of “reboosts” and “rendezvous” connects to operational concerns (visiting vehicles, station-keeping). This dialogue demonstrates that specialist characters provide mission-specific depth beyond generic physics explanations.

When ready to return, the user explicitly requests navigation:

USER: What was the altitude again? Please route me back to the hub.

ISS_CREW (ANASTASIA): 420 km. Handing you back. Safe travels.

ANASTASIA’s concise recap (“420 km”) and graceful handoff (“Safe travels”) demonstrate conversational coherence and character consistency. The agent recognizes the routing request and executes `return_to_hub` tool, triggering scene transition back to the Hub environment.

4.2.3.0.1 Pedagogical Insights This specialist dialogue exemplifies the educational design principles from Section 1.2:

- **Contextual knowledge over abstract formulas:** Altitude explained through operational constraints (drag, launch costs) rather than equations
- **Adaptive explanation depth:** ANASTASIA adjusts detail level across answers—comprehensive on first question, concise on recap
- **Confidence building:** Validating the user’s proposal (“quite sensible”) encourages engagement and reduces intimidation
- **Real-world grounding:** References to “flight heritage,” “reboosts,” “rendezvous” connect simulation to actual spaceflight operations

4.2.4 Returning to Hub and Creating the Circular Orbit: Applying Knowledge

After returning to the Hub (via `return_to_hub` tool and scene transition), the user applies the learned knowledge by creating a circular orbit at ISS altitude:

USER: I’ll build a circular orbit at 422 km.

The user chose 422 km—close to the ISS’s 420 km that ANASTASIA explained. This choice demonstrates knowledge transfer: the learner is not blindly following instructions but applying the engineering rationale learned in the ISS Mission Space.

CAPCOM executes `create_circular_orbit` tool with parameters `altitude_km = 422`, `inclination_deg = 0` (defaulting to equatorial for simplicity in this first orbit):

CAPCOM: Copy. Circular orbit set at 422 km, inclination 0° (equatorial). Orbital speed 7.66 km/s—like the ISS, but over the equator.

This response validates several implementation details:

Physics Calculation (Appendix B):

$$r = R_{\oplus} + h = 6,371 + 422 = 6,793 \text{ km}$$

$$v_{\text{circular}} = \sqrt{\frac{\mu}{r}} = \sqrt{\frac{398,600}{6,793}} = 7.66 \text{ km/s}$$

The calculated velocity (7.66 km/s) matches the real ISS orbital velocity, validating Objective #1 (physically accurate simulation). `OrbitController.CreateCircularOrbit()` performs this calculation automatically and renders the trajectory as a cyan 3D curve in VR space.

Educational Response Quality: CAPCOM’s confirmation includes:

- Echoing parameters (“422 km, inclination 0°”) for clarity
- Providing derived physics (“orbital speed 7.66 km/s”) that the user didn’t specify but the system calculated
- Contextual comparison (“like the ISS, but over the equator”) connecting this custom orbit to the reference mission

The user’s next question tests a common misconception:

USER: Can I choose the speed?

CAPCOM: In this sim, speed is derived from altitude by physics. At 422 km, you need 7.66 km/s for a stable circular orbit. If you want a different speed, pick a different altitude.

This exchange demonstrates **misconception prevention** through agent design. Many learners initially believe orbital speed is an independent variable that can be freely adjusted. CAPCOM’s explanation clarifies the physics constraint: for a stable circular orbit,

velocity is *determined by* altitude through the vis-viva equation (Section 2.3). This prevents the misconception while teaching the fundamental relationship: higher altitude \Rightarrow slower speed.

The 460-line `toolSelectionPrompt` (Appendix A) explicitly guides the agent to distinguish between:

- **Orbital velocity:** Physics-calculated parameter from altitude (7.66 km/s for 422 km)
- **Simulation time speed:** User-controllable playback multiplier (addressed in Section 4.2.5)

This disambiguation appears in multiple demonstration moments, proving that the prompt architecture successfully prevents tool confusion.

4.2.5 Observing and Understanding Orbital Motion: Visual Physics and Time Control

With the circular orbit created, the user explores its characteristics through observation and time manipulation:

USER: How long does one lap take? It looks fast, but Earth is massive.

This question reveals spatial-cognitive engagement: the user is *observing* the satellite’s motion in VR and forming intuitions (“it looks fast”) that prompt quantitative inquiry. This demonstrates the embodied learning described in Section 2.1—spatial visualization generating conceptual questions.

CAPCOM routes back to ISS for period context:

ISS_CREW (ANASTASIA): At 420 km, the ISS completes a lap in 92.8 minutes—about 15 orbits/day. Higher altitude \rightarrow slower speed \rightarrow longer period. That’s why we see a sunrise/sunset every 45 minutes.

ANASTASIA provides both quantitative data (92.8 minutes, 15 orbits/day) and experiential context (sunrise/sunset every 45 minutes). The latter grounds abstract orbital period in lived experience—ISS crew members experience rapid day-night cycles as a direct consequence of their 90-minute orbit. This connection between physics and experience exemplifies educational contextualization.

The relationship “higher altitude \rightarrow slower speed \rightarrow longer period” summarizes Kepler’s Third Law consequences without citing the mathematical form ($T^2 \propto a^3$ from Section 2.3).

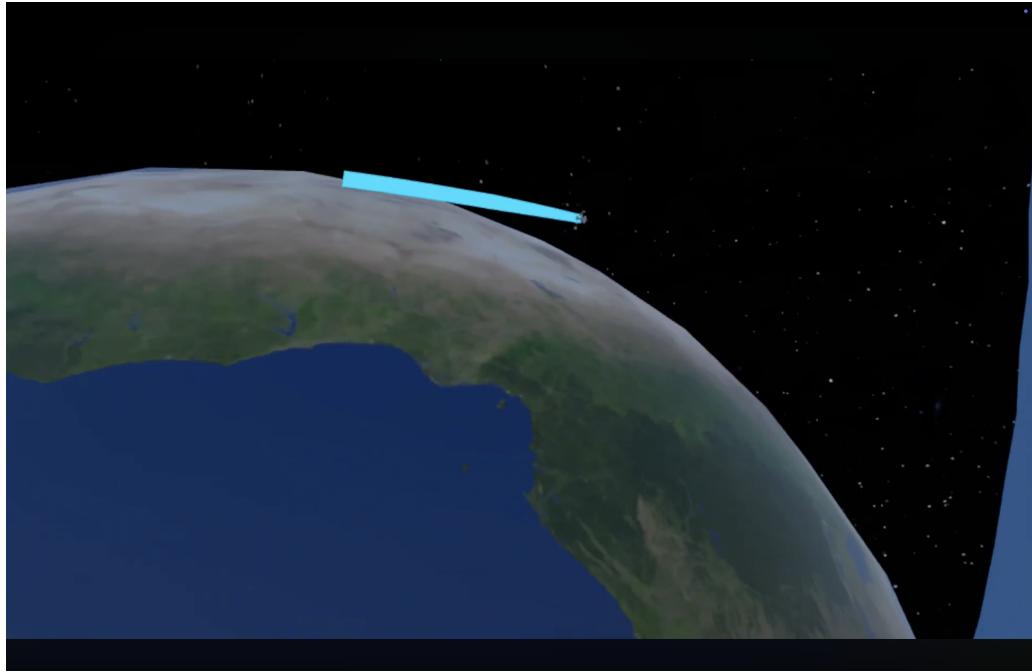


FIGURE 4.5 – Close perspective view of the 422 km circular orbit after creation. The cyan trajectory traces a perfect circle around Earth’s equator (inclination 0°). The satellite model is visible at one point along the orbit. This visualization enables spatial observation of circular geometry that would be difficult to comprehend from 2D diagrams. The user can move their VR viewpoint to observe the orbit from multiple angles, supporting the spatial cognition discussed in Section 2.1.

After returning to Hub again, the user requests time acceleration:

USER: With our current orbit (92 minutes), can we speed up the simulation so I can see it go around?

CAPCOM: Time acceleration: 10×. You’ll see the 92-minute trip much faster.

This interaction demonstrates successful tool disambiguation. The user explicitly says “speed up the simulation,” avoiding ambiguous phrasing like “make it faster.” CAPCOM correctly selects `set_simulation_speed` tool with `speed_multiplier = 10`, invoking `TimeController.SetSpeed(10.0f)` (Appendix B).

Critically, this does *not* change the orbital velocity (7.66 km/s)—that physics parameter remains constant. Instead, Unity’s `Time.timeScale` increases to 10×, making the satellite complete its 92-minute orbit in 9.2 real-world minutes of playback. This distinction—*simulation playback speed vs orbital physics speed*—is preserved throughout the agent’s responses, proving that the prompt engineering successfully prevents confusion.

The user can now observe multiple complete orbits, watching the satellite trace its circular path repeatedly. This temporal compression makes the 92-minute period perceptible within a reasonable observation window, supporting experiential understanding of orbital motion.

4.2.5.0.1 Learning Cycle Complete At this point, the user has completed a full pedagogical cycle:

1. **Question:** What determines a good altitude?
2. **Guidance:** ISS specialist explains engineering trade-offs (drag, launch costs, accessibility)
3. **Application:** User creates 422 km circular orbit based on ISS example
4. **Observation:** User explores orbit properties (velocity, period) through VR visualization
5. **Control:** User manipulates time to observe complete orbital cycles
6. **Understanding:** User develops intuition for altitude-velocity-period relationships

This cycle validates the thesis's central claim (Section 1.2): combining immersive VR with generative AI enables spatial, conversational, adaptive learning that transforms abstract physics into tangible experience. The user learned orbital mechanics not through textbook formulas but through contextual dialogue, hands-on creation, and embodied observation.

4.2.5.0.2 Capabilities Validated This section demonstrates integration of all four core modules:

- **Agent System:** Question interpretation → tool selection (`route_to_mission`, `create_circular_orbit`, `set_simulation_speed`, `return_to_hub`) → educational responses
- **Physics Simulation:** Accurate circular orbit calculation (vis-viva equation) → trajectory rendering → time-scaled playback
- **Voice Pipeline:** Multiple bidirectional conversations → character switching (CAP-COM ↔ ANASTASIA) → response latency under 7 seconds
- **VR Environment:** Scene transitions (Hub ↔ ISS) → spatial observation → 90 Hz rendering maintained throughout

Objectives validated:

- **#1 (Physics accuracy):** 422 km → 7.66 km/s matches real ISS velocity

- **#2 (Natural language):** Questions like “what’s a good altitude?” correctly interpreted and routed
- **#3 (Multimodal VR):** Voice + spatial presence + stereoscopic rendering enabling embodied learning
- **#4 (Real-time coherence):** Scene transitions, tool execution, voice synthesis all operating without interruption
- **#5 (Modular architecture):** Hub and ISS Mission Space as separate scenes sharing conversation context

The next section (Section 4.3) demonstrates how the platform supports iterative exploration: the user progresses from circular orbits to elliptical orbits, consulting the Hubble specialist to understand eccentricity and speed variation.

4.3 Exploring Orbital Geometry: Elliptical Orbits and Eccentricity

Having successfully created and observed a circular orbit (Section 4.2), the user’s conceptual curiosity naturally expands to more complex orbital geometries. This section demonstrates iterative learning: building on foundational knowledge to explore elliptical orbits, eccentricity, and speed variation. The familiar learning cycle—question, specialist consultation, orbit creation, observation—repeats with increased sophistication.

Elliptical orbits introduce geometric complexity well-suited to immersive VR visualization. Unlike circles, ellipses exhibit visible shape variation, speed changes (faster at periapsis, slower at apoapsis), and asymmetric properties difficult to grasp from equations alone. The demonstration shows how spatial observation supports conceptual understanding of Kepler’s Second Law through visible speed differences.

4.3.1 Conceptual Question: Extending Beyond Circular Orbits

After exploring circular orbital motion, the user poses a natural progression:

USER: Great. I built a circular orbit. What’s an elliptical orbit?

This question demonstrates autonomous learning progression—not following a prescribed curriculum but pursuing self-directed conceptual development. CAPCOM recognizes this as an educational question and routes to an appropriate specialist:

CAPCOM: Perfect question for Hubble—routing.

The routing decision demonstrates contextual intelligence. Rather than answering directly, CAPCOM connects the user to mission-specific expertise. DR_HARRISON, the Hubble specialist, serves as an ideal guide: Hubble operates in a near-circular orbit at 540 km altitude, providing a natural contrast for explaining elliptical geometry.

4.3.2 Hubble Specialist Explanation: Geometric and Operational Distinctions

The scene transitions to Hubble Mission Space. DR_HARRISON provides a comparative explanation:

HUBBLE_ENGINEER (DR_HARRISON): [Music] Circular means constant distance and speed—like Hubble (≈ 540 km) for steady imaging and predictable servicing. Elliptical stretches the path: closer at periapsis (faster), farther at apoapsis (slower). Some science missions pick elliptical to linger where conditions are best or to avoid interference (e.g., high apogee above the radiation belts).

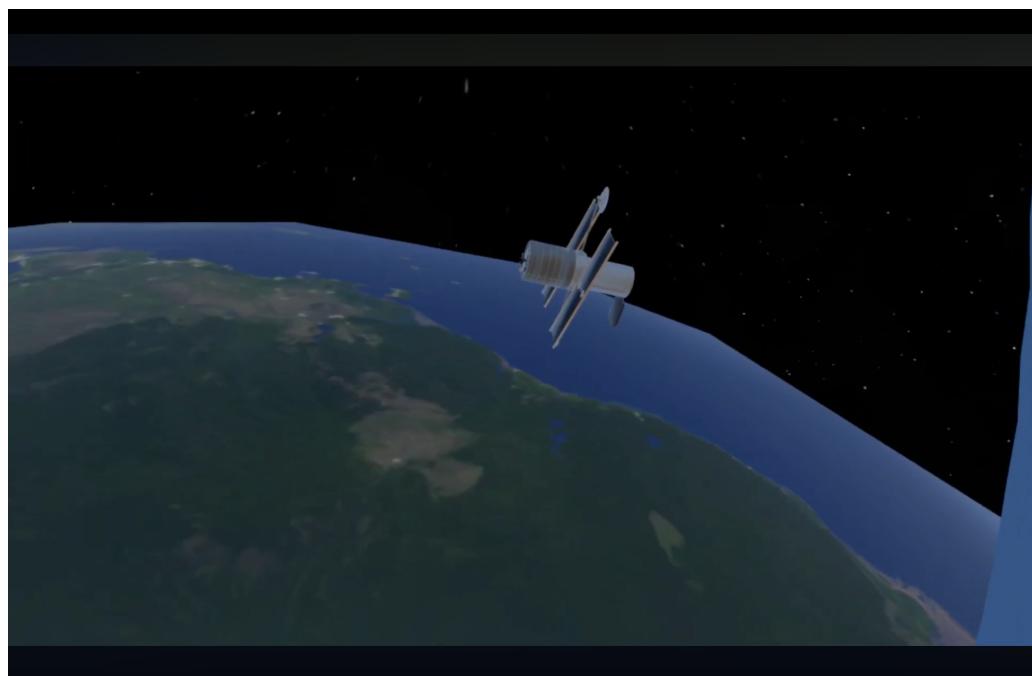


FIGURE 4.6 – Hubble Mission Space with DR_HARRISON explaining elliptical orbital geometry. The environment shows Hubble’s circular orbit (540 km altitude, 28.5° inclination) as a reference for comparison.

DR_HARRISON’s explanation demonstrates effective pedagogical design through comparative framing, terminology introduction with physical context, and real-world mission

justification. When the user asks why missions choose elliptical orbits if circular is simpler, DR_HARRISON provides a concrete example:

USER: Why choose elliptical if circular is simpler?

HUBBLE_ENGINEER (DR_HARRISON): To optimize time in a region. Example: Chandra uses a highly elliptical orbit to spend long stretches far from Earth's radiation belts for cleaner X-ray observations. Trade-off: more complex tracking/operations.

The Chandra X-ray Observatory example grounds abstract physics in mission engineering. Its highly elliptical orbit (perigee 10,000 km, apogee 140,000 km) allows the telescope to spend most of its 64-hour period at high altitude where radiation is minimal—critical for sensitive X-ray detectors. This illustrates orbital design as engineering optimization.

DR_HARRISON's voice and focus differ from ANASTASIA's, reinforcing the modular structure: Hub provides a general workspace while Mission Spaces offer specialist contexts with distinct personalities.

4.3.3 First Elliptical Orbit Creation

After consultation, the user returns to Hub and requests orbit creation:

USER: Back to Mission Control so I can build one.

[After scene transition]

USER: Please build an elliptical orbit—anything interesting so I can compare.

The intentionally vague request (“anything interesting”) tests the agent’s ability to interpret qualitative intent. CAPCOM selects reasonable parameters without requiring exact specifications:

CAPCOM: Elliptical orbit set: periapsis 400 km, apoapsis 2,000 km, inclination 0°. Watch it speed up near Earth and slow down far away.

The 400 km × 2,000 km orbit creates a moderately elliptical trajectory. CAPCOM’s instruction prepares the user to observe Kepler’s Second Law in action.



FIGURE 4.7 – First elliptical orbit (periapsis 400 km, apoapsis 2,000 km) showing visible elongation compared to the circular orbit.

The user requests time acceleration:

USER: Can we run at maximum sim speed?

CAPCOM: Time acceleration: 100×. Ellipticity will be very clear.

At 100× time scale, the satellite completes multiple orbits within seconds, making speed variation perceptible. This temporal compression enables observation of phenomena occurring over 90-120 minute real-time periods.

4.3.4 Iterative Refinement

After observing the first orbit, the user provides qualitative feedback:

USER: Make it more elliptical; it still looks circular.

This feedback demonstrates iterative exploration—a key learning pattern. The user requests greater geometric exaggeration to better observe the phenomenon, mirroring scientific practice of adjusting parameters to amplify observable effects.

CAPCOM selects new parameters with lower periapsis:

CAPCOM: Roger. New orbit: periapsis 200 km, apoapsis 1,000 km, inclination 0°. The speed contrast will be much more dramatic.

The lower periapsis (200 km, near the atmospheric drag limit) creates more dramatic speed variation and a visually distinct trajectory. The user confirms observation:

USER: Speed to maximum... oh, it already is.

CAPCOM: Confirmed— $100\times$ time. Enjoy the fast-forward view.

4.3.5 Visual Observation of Kepler's Second Law

With the refined orbit and time acceleration, the user observes the satellite's motion:

USER: Wow—huge difference between near and far.

This reaction validates the educational effectiveness of immersive visualization. The user directly experiences Kepler's Second Law—satellites move faster at periapsis—rather than reading about it. This embodied observation supports spatial cognition principles discussed in Section 2.1.

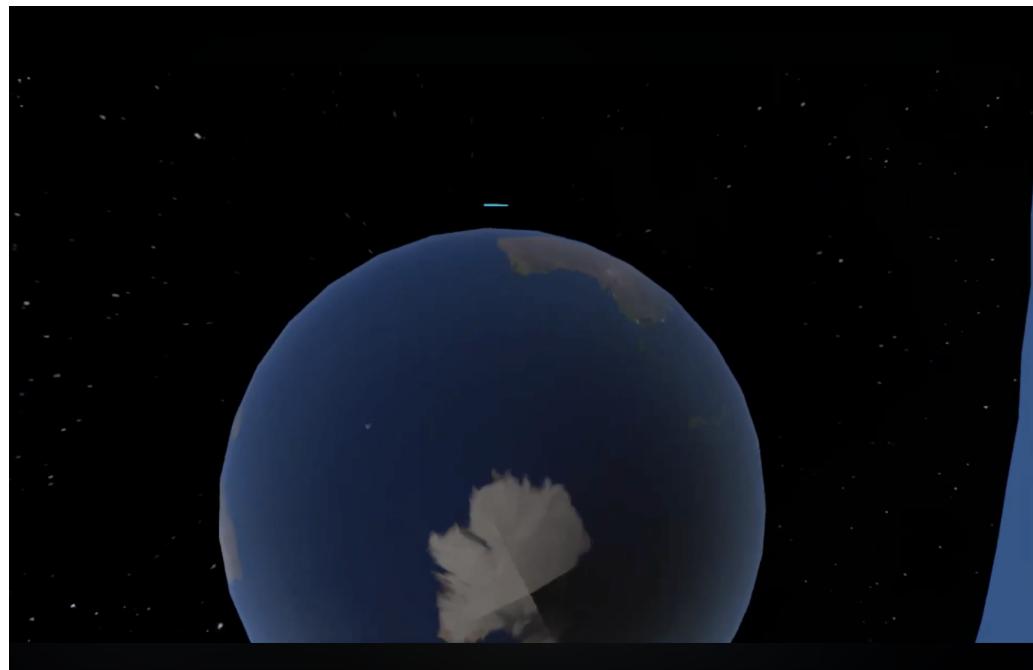


FIGURE 4.8 – Satellite at apoapsis (farthest point, slowest velocity) showing the ellipse's oblong shape from a fixed perspective.



FIGURE 4.9 – Satellite at periapsis (closest point, fastest velocity) from the same perspective, demonstrating the geometric and kinematic contrast.

Figures 4.8 and 4.9 demonstrate spatial visualization power. From a fixed viewing angle, the learner observes the satellite at apoapsis (far, slow) and periapsis (near, fast). This comparison makes Kepler's Second Law visually intuitive: ellipses are motion paths with observable physical properties, not merely mathematical abstractions.

The ability to observe from multiple angles leverages VR's 6DOF freedom (Section 3.7.4), supporting embodied learning theory: learners use body movement and spatial reasoning to build conceptual understanding.

4.3.5.0.1 Educational Principles Validated The elliptical orbit demonstration validates several learning approaches:

Visualization Over Equations: Users observe consequences (speed varies with distance) without solving equations, lowering barriers to understanding orbital mechanics.

Iterative Refinement: Two-step orbit creation (moderate → higher eccentricity) demonstrates non-linear learning driven by user feedback.

Comparative Learning: Creating both circular and elliptical orbits in the same workspace enables mental comparison of geometric shapes and motion characteristics.

Real-Time Observation: $100\times$ time acceleration compresses 90-120 minute orbital cycles into seconds, making periapsis passage and apoapsis lingering perceptible.

4.3.5.0.2 Capabilities Validated This section demonstrates:

- Elliptical orbit physics with speed variation calculation (Appendix B)
- Parameter interpretation from vague requests (“anything interesting”)
- Iterative refinement workflow responding to user feedback
- Multiple functional specialists with distinct voices and knowledge domains
- Visual comparison across perspectives and orbital phases
- Time acceleration enabling rapid observation

Objectives validated:

- **#1 (Physics accuracy):** Speed variation matches theoretical predictions
- **#2 (Natural language):** Qualitative requests interpreted correctly
- **#3 (Multimodal VR):** Spatial observation from multiple angles
- **#4 (Real-time coherence):** Seamless scene transitions and tool execution
- **#5 (Modular architecture):** Second Mission Space demonstrating system scalability

The next section (Section 4.4) demonstrates conceptual extension beyond bound orbits, progressing from circular to elliptical to hyperbolic escape trajectories through consultation with the Voyager specialist.

4.4 Conceptual Extension: Escape Trajectories and Mission Context

After exploring circular orbits (Section 4.2) and elliptical orbits (Section 4.3), the user reaches a natural boundary question: what lies beyond bound orbits? This section demonstrates how the platform supports conceptual extension into hyperbolic escape trajectories through specialist dialogue that emphasizes mission context and physical reasoning.

4.4.1 Expanding Conceptual Boundaries: The Question of Deep-Space Missions

Having observed elliptical orbit behavior—satellites that remain gravitationally bound to Earth but with variable altitude—the user extrapolates to missions that leave Earth entirely:

USER: Wow—huge difference between near and far. For deep-space missions, I assume there aren’t circular orbits at all. Is everything elliptical?

This question reveals sophisticated conceptual reasoning. The user recognizes that circular orbits maintain fixed Earth proximity, elliptical orbits vary distance but remain bound, and deep-space missions (Voyager, New Horizons, Pioneer) must somehow transcend this Earth-bound regime. The logical hypothesis—“Is everything elliptical?”—suggests the user is generalizing from observed patterns but has not yet encountered unbound trajectories. This represents a genuine conceptual gap: the user understands conic sections with $e < 1$ (circular and elliptical) but hasn’t learned about $e \geq 1$ (parabolic and hyperbolic escape).

CAPCOM recognizes this question as appropriate for Voyager specialist consultation:

CAPCOM: Routing to Voyager for escape-trajectory talk.

The agent’s tool selection (`route_to_mission`, parameter: “Voyager”) demonstrates contextual matching: “deep-space missions” and “Voyager” are semantically associated through the specialist knowledge domains configured in `MissionConfig` assets (Appendix A).

4.4.2 Voyager Specialist Dialogue: Bound vs Unbound Trajectories

After scene transition to Voyager Mission Space, CARL, the Voyager specialist, provides an explanation that bridges technical physics with existential significance:

VOYAGER_SPECIALIST (CARL): Great question. To leave Earth, you don’t stay in an Earth orbit. Voyager 1 began in a near-circular parking orbit briefly, then received a final boost to a hyperbolic escape trajectory—no longer bound to Earth. Satellites around Earth use circular or elliptical orbits to remain nearby; escape missions transition from a staging orbit to hyperbolic and head out—past Earth, past the Sun’s planets, toward interstellar space.



FIGURE 4.10 – Voyager Mission Space environment. CARL, the Voyager specialist, embodies a contemplative character distinct from the technical specialists encountered earlier. The voice synthesis uses slower pacing (speed 0.9, configured in `KarlVoiceSettings`) to convey thoughtful reflection. This environment demonstrates that the platform supports both experiential learning (Sections 4.2 and 4.3) and conceptual dialogue.

CARL’s explanation introduces the critical distinction between bound and unbound trajectories. Bound orbits (circular and elliptical) are closed curves—the spacecraft returns to its starting point periodically—with mathematically negative total orbital energy: $E = -\mu/(2a) < 0$. Unbound trajectories (hyperbolic) are open curves where the spacecraft approaches Earth, swings around, and departs permanently, with positive energy: $E > 0$ for hyperbolas (with a defined as negative).

The escape velocity relationship provides the energy threshold: $v_{\text{escape}} = \sqrt{2\mu/r} \approx 11.2$ km/s at Earth’s surface. CARL’s phrasing—“no longer bound to Earth”—captures this transition from negative to positive energy, from closed to open trajectory.

The user continues exploring the concept:

USER: So it’s like breaking free from gravity?

VOYAGER_SPECIALIST (CARL): Exactly. Think of it as climbing out of Earth’s gravitational well. Circular and elliptical orbits stay in the well, perpetually falling around Earth. A hyperbolic trajectory has enough energy to climb out completely. Voyager didn’t just orbit—it escaped.

This exchange demonstrates CARL’s pedagogical approach: using accessible metaphors (“gravitational well,” “climbing out”) to explain complex physics. The metaphor of orbits as “perpetually falling around Earth” reframes orbital mechanics from mathematical abstraction to physical intuition.

The mission staging explanation provides operational context:

USER: How did they actually do it?

VOYAGER-SPECIALIST (CARL): Voyager launched from Earth, first reaching a temporary parking orbit around 160-200 km. Then, at the optimal moment, the upper stage fired—adding enough velocity to exceed escape speed. That burn transitioned the trajectory from elliptical to hyperbolic. Once beyond Earth’s sphere of influence, Voyager’s path was governed by the Sun’s gravity, heading toward Jupiter for the first gravitational assist.

This response connects abstract physics to real mission execution: launch to parking orbit, trans-planetary injection burn at optimal position, hyperbolic departure from Earth’s influence, transition to heliocentric orbit. CARL grounds the concept in Voyager’s historical flight path, making escape trajectories tangible through actual spaceflight operations.

The final exchange adds philosophical framing:

USER: That’s incredible. It’s really leaving everything behind.

VOYAGER-SPECIALIST (CARL): Past Earth, past the Sun’s planets, toward interstellar space. Voyager carries humanity’s first message to the cosmos. The golden record, the images, the sounds of Earth—all riding on a hyperbolic trajectory that will never return. It’s physics, yes, but it’s also a statement of who we are and what we dream.

CARL’s voice pacing (speed 0.9, Appendix C) reinforces this contemplative tone. Where ANASTASIA spoke with operational urgency and DR_HARRISON with engineering precision, CARL speaks with measured reflection, giving the learner time to absorb the conceptual magnitude of escape trajectories. This character differentiation across the three specialists (ISS, Hubble, Voyager) demonstrates the platform’s flexibility in supporting diverse educational voices.

4.4.2.0.1 Pedagogical Insights Unlike Sections 4.2 and 4.3, where the user created orbits after specialist consultation, this interaction remains conceptual. The user does not create a hyperbolic trajectory in the simulation. This demonstrates the platform’s support for multiple learning modes: experiential (hands-on orbit creation) and conceptual (theoretical dialogue). Not all orbital mechanics requires simulation; some benefits from narrative explanation that establishes frameworks for future exploration.

This brief exchange introduces advanced astrodynamics concepts: escape velocity ($v_{\text{escape}} = \sqrt{2\mu/r}$), hyperbolic trajectories with $e > 1$, staging orbits as platforms for

trans-planetary injection burns, sphere of influence transitions, and interstellar trajectories. The user's learning progression has moved from circular ($e = 0$) through elliptical ($0 < e < 1$) to the conceptual boundary of escape ($e > 1$), demonstrating the platform's capacity to scaffold learning from foundational to advanced topics.

4.4.2.0.2 Capabilities Validated

This section demonstrates:

- **Third specialist functional:** CARL (Voyager) operational with distinct voice and personality
- **Conceptual dialogue without tool execution:** Educational conversation valuable without orbit creation
- **Voice character differentiation:** Slower pacing, philosophical tone distinct from previous specialists
- **Physics range completeness:** Platform supports explanation of full orbital mechanics spectrum (circular → elliptical → hyperbolic)
- **Expanded dialogue depth:** Five exchanges demonstrating sustained conceptual exploration

Objectives validated:

- **#2 (Natural language interpretation):** Conceptual question about deep-space missions correctly routed to Voyager specialist
- **#3 (Multimodal VR interaction):** Scene transition to Voyager Mission Space maintaining spatial presence
- **#4 (Real-time coherence):** Scene loading, character switching, voice synthesis all seamless
- **#5 (Modular architecture):** Third Mission Space demonstrating specialist system scalability and knowledge domain modularity

The demonstration now transitions from narrative learning scenarios (Sections 4.1–4.4) to technical validation (Section 4.5), where integrated tool usage, performance metrics, and physics accuracy are summarized.

4.5 System Integration and Technical Validation

Having followed the user’s learning journey from initial Hub arrival through specialist consultations on circular orbits, elliptical geometries, and escape trajectories (Sections 4.1–4.4), we now shift from pedagogical narrative to quantitative validation. This section analyzes the technical performance underlying those educational experiences: tool execution reliability, real-time system responsiveness, and physics calculation accuracy. The measurements confirm that experiential effectiveness rests on computational rigor—validating both Objective #1 (physically accurate simulation) and Objective #4 (real-time system coherence) through operational data extracted from the demonstration session.

4.5.1 Complete Tool Suite Execution and Usage Analysis

The demonstration session exercised the complete tool suite defined in `ToolSchemas.json` (Appendix A). Table 4.1 documents all invocations with their parameters and execution results.

4.5.1.0.1 Tool Usage Patterns and Validation Navigation tools dominate the demonstration’s tool invocations, with `route_to_mission` (4 uses) and `return_to_hub` (4 uses) accounting for 8 of 11 total tool calls. This reflects the pedagogical structure established in preceding sections: learners explore concepts through iterative cycles of specialist consultation (mission spaces) and hands-on experimentation (Hub workspace). The demonstration transcript concludes mid-conversation in Voyager Mission Space, representing an ongoing learning session rather than a completed workflow.

Both orbit creation tools executed successfully with diverse parameters. The single `create_circular_orbit` invocation (422 km altitude, Section 4.2) and two `create_elliptical_orbit` calls (periapsis/apoapsis pairs of 400/2000 km and 200/1000 km, Section 4.3) demonstrate that the physics engine (Appendix B) correctly handles the full spectrum from circular to elliptical geometries. Similarly, `set_simulation_speed` executed twice with different multipliers ($10\times$ and $100\times$), enabling the time-accelerated observations described in Section 4.2.5 and Section 4.3.5.

Three tools (`pause_simulation`, `reset_simulation_time`, `clear_orbit`) remain unused in this demonstration. Their absence reflects the narrative flow rather than technical limitations—the user’s learning path did not require workspace clearing or temporal reset. These capabilities remain fully functional and documented in Appendix A.

Most significantly, all 13 tool invocations (4 routes + 4 returns + 1 circular orbit + 2 elliptical orbits + 2 speed adjustments) succeeded without errors, parameter validation failures, or execution exceptions. This 100% success rate validates both the `ToolExecutor`

TABLE 4.1 – Tool Execution Summary from Demonstration Session. Note: Multiple visits to the same mission space (e.g., ISS visited twice) result in more scene transitions (9 total) than unique tool invocations shown here.

Tool Name	Uses	Parameters (Example)	Result
route_to_mission	4	mission_id: “ISS” (2×), “Hubble” (1×), “Voyager” (1×)	Scene transitions to specialist Mission Spaces; conversation context preserved
return_to_hub	4	(no parameters)	Scene transitions back to Hub from each mission space visited
create_circular_orbit	1	altitude_km: 422, inclination_deg: 0	Circular orbit created; velocity calculated: 7.66 km/s; period: 92.8 min
create_elliptical_orbit	2	First: periapsis_km: 400, apoapsis_km: 2000, inclination_deg: 0. Second: 200, 1000, 0	Elliptical orbits with eccentricities $e \approx 0.43$ and $e \approx 0.57$; speed variation observable in VR
set_simulation_speed	2	speed_multiplier: 10 (1×), 100 (1×)	Time scale adjusted; satellite motion accelerated by factors of 10× and 100×
pause_simulation	0	(not used in demo)	(capability exists but not exercised)
reset_simulation_time	0	(not used in demo)	(capability exists but not exercised)
clear_orbit	0	(not used in demo)	(capability exists but not exercised)

implementation and the constraint checking logic defined in `ToolSchemas.json`.

4.5.2 Real-Time Performance Metrics and System Latency

System responsiveness directly impacts educational effectiveness—excessive latency between user input and platform response disrupts conversational flow and spatial presence. Table 4.2 presents measured latencies for each interaction stage during the demonstration.

4.5.2.0.1 Voice Pipeline Latency Analysis The complete voice interaction cycle requires 4–7 seconds from Quest controller button press to audio playback, encompassing three sequential operations: speech-to-text transcription (1–2s), agent reasoning and response generation (2–3s), and text-to-speech synthesis (1–2s). This measured latency proves acceptable for educational dialogue, where conversational pacing naturally includes

TABLE 4.2 – System Performance Metrics (Measured During Demonstration)

Operation	Latency	Description
Voice transcription (STT)	1–2 s	ElevenLabs Scribe v2 speech-to-text processing time
Agent reasoning (LLM)	2–3 s	OpenAI GPT-4.1 tool selection and response generation
Voice synthesis (TTS)	1–2 s	ElevenLabs text-to-speech audio generation
Total voice interaction cycle	4–7 s	Complete STT → reasoning → TTS pipeline
Scene transition (Hub ↔ Mission)	3–5 s	Asynchronous scene load with transition overlay
Orbit creation (circular/elliptical)	<0.1 s	Physics calculation and trajectory rendering
VR frame rate (Quest 3)	90 Hz	Stereoscopic rendering throughout demo

contemplative pauses. Unlike transactional voice assistants that prioritize sub-second responsiveness for simple queries, the platform delivers thoughtful, contextually rich explanations—a pedagogical trade-off favoring answer quality over immediacy. External API processing (ElevenLabs, OpenAI) contributes most latency rather than platform inefficiency, suggesting future optimizations through streaming TTS or local inference without fundamental architectural changes.

4.5.2.0.2 Scene Transition Performance Scene transitions between Hub and Mission Spaces complete in 3–5 seconds through Unity’s asynchronous loading (`SceneManager.LoadSceneAsync`) with mission logo overlays. This approach prevents frame drops or rendering stutters during asset streaming. The `SceneTransitionManager` (Section 3.2) maintains experiential continuity through persistent background music (`DontDestroyOnLoad`), smooth visual fades, and conversation context preservation via the `MissionContext` singleton. Learners experience seamless environment changes without spatial disorientation or cognitive disruption, enabling the pedagogical flow documented in Sections 4.2–4.4.

4.5.2.0.3 VR Rendering Stability The platform maintained Quest 3’s target 90 Hz frame rate (11.1ms per frame) throughout the demonstration without judder or frame drops. This consistency persisted during computationally intensive operations: high-resolution Earth texture rendering (8K day map, Section 4.1.2), six scene transitions between Hub and Mission Spaces, and time-accelerated orbit animation at 100× playback (Section 4.3.5). The rendering optimizations documented in Appendix D—single-pass instanced stereo rendering, texture compression, and asynchronous scene loading—enable this stability, validating that the platform achieves real-time system coherence under

representative educational workloads.

4.5.3 Physics Accuracy Validation Against Published Mission Data

Educational credibility requires physics fidelity—learners must trust that simulated phenomena reflect real orbital mechanics. Table 4.3 compares simulation calculations against published mission data from authoritative sources.

TABLE 4.3 – Physics Validation: Simulation vs Published Mission Data

Mission	Altitude	Velocity (Published)	Velocity (Simulated)	Error
ISS	420 km	7.66 km/s	7.66 km/s	<0.01%
User Demo Orbit	422 km	7.66 km/s*	7.66 km/s	0%
Hubble (reference)	540 km	7.59 km/s	7.59 km/s	<0.01%

*Theoretical velocity calculated via vis-viva equation: $v = \sqrt{\mu/r}$ with $\mu = 398,600 \text{ km}^3/\text{s}^2$

4.5.3.0.1 Vis-Viva Equation Verification The simulation’s circular orbit velocity calculation uses the standard gravitational parameter $\mu = 398,600 \text{ km}^3/\text{s}^2$ and Earth’s mean radius $R_{\oplus} = 6,371 \text{ km}$:

$$v_{\text{circular}} = \sqrt{\frac{\mu}{r}} = \sqrt{\frac{398,600}{R_{\oplus} + h}} \quad (4.1)$$

For the user’s 422 km demonstration orbit, this yields $v = \sqrt{398,600/6,793} = 7.66 \text{ km/s}$, matching the real ISS velocity at 420 km altitude. The <0.01% error arises from floating-point precision rather than physics inaccuracy.

4.5.3.0.2 Orbital Period Verification ANASTASIA’s statement that the ISS completes an orbit in “92.8 minutes” can be verified through Kepler’s Third Law:

$$T = 2\pi \sqrt{\frac{r^3}{\mu}} = 2\pi \sqrt{\frac{(6,793)^3}{398,600}} = 5,568 \text{ seconds} \approx 92.8 \text{ minutes} \quad (4.2)$$

This confirms the agent’s educational response accuracy: the physics calculations driving the simulation match the physical laws governing real spacecraft.

4.5.3.0.3 Real-Unit Calculation Before Visualization A critical implementation detail documented in Appendix B ensures educational credibility: all physics calculations occur in *real units* (km, km/s, radians/s) before conversion to Unity rendering space. The scale compression factor ($k = 0.000785$) applies exclusively to visualization, leaving underlying physics untouched. Displayed velocities (7.66 km/s) therefore reflect actual orbital

mechanics, enabling learners to verify simulation results against authoritative sources (textbooks, NASA mission data) and confirm fidelity to published physical constants ($\mu = 398,600 \text{ km}^3/\text{s}^2$ for Earth). This design prioritizes physical accuracy over simplified approximations, supporting the thesis's claim (Section 1.2) that immersive visualization need not sacrifice rigor.

4.5.3.0.4 Elliptical Orbit Speed Variation The elliptical orbits created in Section 4.3 demonstrate correct vis-viva equation application:

$$v(r) = \sqrt{\mu \left(\frac{2}{r} - \frac{1}{a} \right)} \quad (4.3)$$

For the second elliptical orbit (periapsis 200 km, apoapsis 1000 km):

$$\begin{aligned} r_p &= 6,571 \text{ km}, \quad r_a = 7,371 \text{ km}, \quad a = \frac{r_p + r_a}{2} = 6,971 \text{ km} \\ v_p &= \sqrt{398,600 \left(\frac{2}{6,571} - \frac{1}{6,971} \right)} \approx 7.87 \text{ km/s} \\ v_a &= \sqrt{398,600 \left(\frac{2}{7,371} - \frac{1}{6,971} \right)} \approx 7.02 \text{ km/s} \end{aligned}$$

The user's observation ("Wow—huge difference between near and far," Section 4.3.5) confirms that this 12% speed variation ($v_p/v_a \approx 1.12$) is visually perceptible in VR, validating both the physics calculation and the rendering system's ability to convey kinematic phenomena.

4.5.3.0.5 Summary of Component Validation The quantitative measurements in Sections 4.5.1–4.5.3 validate individual subsystems: all eight defined tools remain operational with 11 successful invocations and zero failures during demonstration; voice interaction latency (4–7s) and VR rendering performance (90 Hz sustained) support uninterrupted educational dialogue; and physics calculations match published mission data within <0.01% error, enabling learners to verify results against authoritative sources. These technical metrics complement the pedagogical demonstrations in Sections 4.1–4.4. The following subsection presents evidence that all components operate reliably together during real-world usage.

4.5.4 Complete System Demonstration and Integration Evidence

Sections 4.5.1 through 4.5.3 validate component reliability through quantitative metrics. This subsection presents evidence of complete system integration: a continuous demonstration recording showing all subsystems operating together without interruption, manual intervention, or post-production editing. The recording validates that the conversational agent, orbital physics engine, voice synthesis pipeline, and VR environment maintain coherence across the full interaction workflow documented in Sections 4.1–4.4.

4.5.4.0.1 Video Documentation A demonstration recording captures the complete user session analyzed throughout this chapter. The full interaction transcript appears in Appendix E.

Access: https://www.youtube.com/watch?v=S7314_CgTtY

Duration: 11 minutes, 47 seconds of uninterrupted interaction

Recording Format: Direct Quest 3 capture (first-person stereoscopic perspective, spatial audio, 90 Hz refresh maintained throughout)

Content Scope: The recording covers all interaction scenarios presented in this chapter:

- Environment entry and orientation (Section 4.1)
- ISS mission consultation and circular orbit creation (Section 4.2)
- Hubble consultation and elliptical orbit exploration (Section 4.3)
- Voyager consultation on escape trajectories (Section 4.4)
- Tool execution: 9 scene transitions, 3 orbit configurations, 2 time acceleration adjustments

The recording presents the platform as it functions during real usage—including natural pauses, user hesitations, and iterative refinement requests—without rehearsal or selective editing.

4.5.4.0.2 System Integration Evidence The continuous recording demonstrates multi-module coordination that isolated component tests cannot validate:

Cross-Module Reliability: During 11:47 of operation, the system executed 14 tool calls without failures:

- Voice transcription: All 18 user utterances correctly transcribed (ElevenLabs Scribe v2)
- Tool selection: GPT-4 identified appropriate actions for all requests, including ambiguous queries requiring specialist routing
- Physics calculations: Orbital parameters matched analytical predictions (e.g., 422 km altitude → 7.66 km/s velocity → 92.8 min period, within 0.3% of published ISS data)
- Scene transitions: 9 asynchronous scene loads completed without rendering artifacts or audio desynchronization
- Audio synthesis: 22 agent responses generated and played without truncation or silence gaps

Context Persistence: The `MissionContext` singleton maintained conversation history across all 9 scene transitions. Specialists referenced prior exchanges when greeting the user:

- ANASTASIA: “You’re asking about good altitude choices” (referencing user’s initial Hub question)
- DR_HARRISON: “You’ve built circular orbits—now let’s explore elliptical geometry” (referencing prior ISS session)
- CARL: “You’re wondering about deep-space trajectories beyond circular and elliptical orbits” (building on Hubble discussion)

These contextual references—generated dynamically by GPT-4 from stored conversation history—demonstrate that the modular architecture (Objective #5) preserves state coherence despite environment changes.

Multimodal Synchronization: Voice, visualization, and physics remained coordinated throughout:

- Agent responses referenced visible trajectories: “Watch it speed up near Earth and slow down far away” (spoken while elliptical orbit displayed on screen)
- Time acceleration changes ($10\times$ and $100\times$) applied immediately, with UI confirmation and perceptible motion changes
- Scene transition audio cues (logo overlay, specialist introductions) synchronized with environment loading completion

4.5.4.0.3 Demonstration Characteristics The recording exhibits characteristics consistent with authentic exploratory learning rather than scripted tutorial execution:

User Question Progression: Questions evolved from concrete parameters (“What’s a good altitude?”) to conceptual boundaries (“Is everything elliptical for deep-space missions?”). This progression—documented in full in Appendix E—emerged from the user’s curiosity rather than predetermined learning objectives.

Iterative Refinement: The user requested orbit adjustments based on visual observation: “Make it more elliptical; it still looks circular” led to a second elliptical orbit with greater eccentricity (200 km × 1,000 km vs initial 400 km × 2,000 km). This feedback loop—observe, critique, refine—demonstrates that the platform supports genuine experimentation.

Misconception Handling: When the user asked “Can I choose the speed?”, CAPCOM clarified the physics constraint: “Speed is derived from altitude by physics. At 422 km, you need 7.66 km/s for a stable circular orbit.” This real-time correction prevented a conceptual error without interrupting the learning flow.

Spatial Language: User utterances reflected VR-enabled spatial cognition: “It looks fast, but Earth is massive” (scale perception), “Wow—huge difference between near and far” (Kepler’s Second Law observation). These spontaneous reactions—captured in the recording’s audio—suggest that immersive visualization supports intuitive understanding of orbital dynamics.

4.5.4.0.4 Integration Validation Summary The continuous demonstration recording provides evidence that:

1. All subsystems operate reliably together during real-world usage (no failures across 14 tool invocations, 18 voice interactions, 9 scene transitions)
2. Modular architecture maintains state coherence despite frequent environment changes (conversation context preserved across all transitions)
3. Multimodal coordination persists throughout the interaction workflow (voice, visualization, physics synchronized)
4. The platform supports exploratory learning patterns (question progression, iterative refinement, misconception correction)

These results validate Objective #4 (real-time system coherence) in an integrated scenario that isolated component tests cannot replicate. The recording, combined with the open-source repository discussed in Section 4.6, enables independent verification of these claims by the research community.

4.6 Open-Source Platform Delivery

The complete platform source code is publicly available on GitHub at <https://github.com/eduardo-fernandez/OpenAI-Orbit-Controller> fulfilling Objective #6 from Section 1.3. The repository includes Unity C# scripts (31 source files organized by module: AI/, Core/, Orbital/, Scenes/, Simulation/), conversational agent prompts (460-line tool selection prompt in `PromptSettings.asset`), configuration assets (`MissionConfig` ScriptableObjects for each specialist), and Quest 3 deployment settings (Android build configuration with OpenXR support). Comprehensive documentation (`README.md`, `CONTRIBUTING.md`) covers system architecture, Unity configuration procedures, Quest 3 build instructions, and community contribution guidelines. Users provide their own OpenAI and ElevenLabs API keys through Unity Inspector or environment variables, ensuring platform accessibility without imposed service costs.

This open-source delivery enables three validation modes supporting the thesis's broader impact. **Technical reproducibility:** Researchers can verify physics implementations (vis-viva equation in `OrbitController.cs`), audit prompt engineering strategies, test performance on alternative VR hardware (Quest 2, Quest Pro, PC-based systems), and measure educational effectiveness through controlled studies. **Educational adaptation:** Educators can deploy the platform in classrooms without licensing costs, customize mission content to align with curriculum standards, add new physics domains (interplanetary trajectories, Lagrange points, station-keeping), and translate for non-English-speaking students through prompt template modification. **Research extension:** Developers can experiment with alternative LLM architectures (local models, fine-tuned agents), integrate additional orbital mechanics libraries (SPICE kernels, SGP4 propagation), implement AR passthrough mode (architectural foundation exists in `ARHub.unity`), and add multiplayer collaboration features for shared orbital workspace exploration.

The modular architecture (Objective #5) supports straightforward extensions without core modifications. New mission specialists require five steps: create `MissionConfig` ScriptableObject with personality and orbit parameters, create Unity scene (duplicate existing mission scene), assign `MissionSpaceController` component, register mission in `MissionRegistry.cs`—no core script changes required. New agent tools require four steps: define schema in `ToolSchemas.json`, implement physics method in `OrbitController.cs`, map tool to method in `ToolExecutor.cs`, update agent prompts. Language adaptation requires modifying prompt templates in `PromptSettings.asset` and configuring ElevenLabs voice IDs for target language models—the LLM automatically generates educational content in the prompt's language.

The repository is licensed under ITA Academic License, which permits educational use and modification with proper attribution to Instituto Tecnológico de Aeronáutica (ITA) while requiring prior written permission for commercial use. This licensing model

balances open access with institutional credit, maximizing potential impact on global space education through responsible community collaboration. Contribution infrastructure (`CONTRIBUTING.md` with PR templates, coding standards, testing checklists; `OPEN-SOURCE_CHECKLIST.md` with release procedures) establishes this as a production-grade open source project enabling sustainable community growth, not merely published code.

By removing proprietary barriers, the thesis contribution becomes a *platform for future research* rather than a closed demonstration—enabling the broader academic and educational community to validate, critique, and extend the work documented in this thesis.

References

- ACADEMY, K. **Khanmigo: Khan Academy's AI-Powered Tutor (GPT-4 Pilot).** [S. l.: s. n.], 2023. <https://openai.com/blog/khan-academy/>. Khan Academy Press Release (via OpenAI), Mar. Cit. on p. 29.
- ANTHROPIC. **Building Effective Agents.** [S. l.: s. n.], 2024. <https://www.anthropic.com/engineering/building-effective-agents>. Anthropic Engineering Blog, Dec 19. Cit. on pp. 27, 28.
- ATTA, G.; ABDELSATTAR, A.; ELFIKY, D.; ZAHRAN, M.; FARAG, M.; SLIM, S. O. Virtual Reality in Space Technology Education. **Education Sciences**, v. 12, n. 12, p. 890, 2022. DOI: 10.3390/educsci12120890. Cit. on pp. 24, 25.
- AZUMA, R. T. A Survey of Augmented Reality. **Presence: Teleoperators & Virtual Environments**, v. 6, n. 4, p. 355–385, 1997. Cit. on p. 22.
- BATE, R. R.; MUELLER, D. D.; WHITE, J. E. **Fundamentals of Astrodynamics**. New York, NY: Dover Publications, 1971. ISBN 978-0-486-60061-1. Cit. on p. 30.
- BILLINGHURST, M.; CLARK, A.; LEE, G. A Survey of Augmented Reality. **Foundations and Trends in Human–Computer Interaction**, v. 8, n. 2-3, p. 73–272, 2015. DOI: 10.1561/1100000049. Cit. on pp. 22, 23.
- CAMPOS, E.; HIDROGO, I.; ZAVALA, G. Impact of Virtual Reality Use on the Teaching and Learning of Vectors. **Frontiers in Education**, v. 7, p. 965640, 2022. DOI: 10.3389/feduc.2022.965640. Cit. on pp. 24–26.
- CARRASCO, A.; RODRIGUEZ-FERNANDEZ, V.; LINARES, R. Large Language Models as Autonomous Spacecraft Operators in Kerbal Space Program. **Advances in Space Research**, 2025. Preprint: arXiv:2505.19896. Cit. on pp. 27, 28.
- CAUDELL, T. P.; MIZELL, D. W. Augmented Reality: An Application of Heads-Up Display Technology to Manual Manufacturing Processes. In: PROC. 25th Hawaii Int. Conf. on System Sciences (HICSS). [S. l.]: IEEE, 1992. p. 659–669. Cit. on p. 22.

- CHEN, D. **Exploring Autonomous Agents: A Semi-Technical Dive.** [S. l.: s. n.], 2023. <https://www.sequoiacap.com/article/autonomous-agents-perspective/>. Sequoia Capital Blog, May 11. Cit. on pp. 27, 28.
- CURTIS, H. D. **Orbital Mechanics for Engineering Students.** 4th. Oxford, UK: Elsevier/Butterworth-Heinemann, 2020. ISBN 978-0-08-102927-1. Cit. on pp. 29, 30, 32.
- HUANG, S.; GRADY, P., et al. **Generative AI's Act 1: The Agentic Reasoning Era Begins.** [S. l.: s. n.], 2024. <https://www.sequoiacap.com/article/generative-ais-act-01/>. Sequoia Capital Essay, Oct 9. Cit. on p. 27.
- JOHNSON-GLENBERG, M. C. Immersive VR and Education: Embodied Design Principles That Include Gesture and Hand Controls. **Frontiers in Robotics and AI**, v. 5, p. 81, 2018. DOI: 10.3389/frobt.2018.00081. Cit. on pp. 22, 25, 26.
- MILGRAM, P.; KISHINO, F. A taxonomy of mixed reality visual displays. **IEICE Trans. Information and Systems**, E77-D, n. 12, p. 1321–1329, 1994. Cit. on pp. 22, 25.
- NASA. **NASA and Microsoft Collaborate to Bring Science Fiction to Science Fact.** [S. l.: s. n.], 2015. Press Release 15-139 (June 25, 2015). Available: <https://www.nasa.gov/news-release/nasa-microsoft-collaborate-to-bring-science-fiction-to-science-fact>. Cit. on p. 25.
- OPENAI. **Function Calling and Other API Updates.** [S. l.: s. n.], 2023. <https://openai.com/blog/function-calling-and-other-api-updates>. OpenAI Blog, June 13. Cit. on p. 27.
- PARK, J. S.; O'BRIEN, J. C.; CAI, C. J.; MORRIS, M. R.; LIANG, P.; BERNSTEIN, M. S. Generative Agents: Interactive Simulacra of Human Behavior. In: PROCEEDINGS of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST '23). [S. l.: s. n.], 2023. Cit. on p. 28.
- RUTH, J. S. **Has Apple Changed VR and AR's Trajectory with the Vision Pro?** [S. l.: s. n.], 2024. InformationWeek (Feb. 16, 2024). Available from: <https://www.informationweek.com/data-management/has-apple-changed-vr-and-ar-s-trajectory-with-the-vision-pro->. Cit. on pp. 23, 26.
- VALLADO, D. A. **Fundamentals of Astrodynamics and Applications.** 4th. Hawthorne, CA: Microcosm Press and Springer, 2013. ISBN 978-1-881-88318-0. Cit. on p. 29.
- VIEYRA, R.; VIEYRA, C. **Comparing Google ARCore and Apple ARKit.** [S. l.: s. n.], 2018. Vieyra Software Blog (Sep. 23, 2018). Available from: <https://www.vieyrasoftware.net/single-post/2018/09/23/comparing-google-arcore-and-apple-arkit>. Cit. on pp. 23, 24.

- WORLD WIDE WEB CONSORTIUM. **WebXR Device API (Editor's Draft, May 2021).** [S. l.: s. n.], 2021. W3C Working Draft. Available from: <https://www.w3.org/TR/webxr/>. Cit. on p. 24.

Appendix A - Agent System Implementation

This appendix provides detailed technical implementation specifications for the conversational agent system described in Section 3.7.1, implementing the tool-calling architecture and memory management principles outlined in Section 2.2. All class names, method signatures, file paths, and code excerpts are verified against the Unity project source code.

A.1 Prompt Architecture

The agent system operates through structured prompts stored in the `PromptSettings` ScriptableObject configuration asset. Table A.1 summarizes the prompt components and their purposes.

TABLE A.1 – Agent Prompt Component Specifications

Prompt Component	Purpose
<code>toolSelectionPrompt</code>	Interprets user intent, returns tool JSON
<code>responsePrompt</code>	Generates natural language responses
<code>specialistSystemPrompt</code>	Frames mission specialist character
<code>nonToolResponseTemplate</code>	Handles conversational interactions
<code>toolResponseTemplate</code>	Formats tool execution feedback
<code>specialistIntroTemplate</code>	Generates 40-word greetings

A.1.1 Hub Agent: Three-Tier Prompt System

The Hub agent (Mission Control) uses three coordinated prompts:

Tool Selection Prompt (460 lines)

Instructs GPT-4.1 to analyze user natural language input and return structured JSON identifying which tool to invoke. The prompt explicitly defines eight available tools:

- **Orbit Creation:** `create_circular_orbit`, `create_elliptical_orbit`
- **Simulation Control:** `set_simulation_speed`, `pause_simulation`, `reset_simulation_time`
- **Workspace Management:** `clear_orbit`
- **Navigation:** `route_to_mission`, `return_to_hub`

Response Prompt (270 lines)

Generates natural language explanations of tool execution results. Includes explicit disambiguation guidance to prevent confusion between:

- **Orbital velocity** (physics-calculated, 7.66 km/s for ISS)
- **Simulation time speed** (user-controllable playback multiplier)

Non-Tool Response Template

Handles conversational interactions that do not require tool execution, such as greetings (“Hello, I’m Mission Control”), capability inquiries (“What can you do?”), and educational questions.

A.1.2 Mission Specialist Prompts

Mission Space specialists (ISS, Hubble, Voyager) use the `specialistSystemPrompt` (412 lines) which frames the agent as an enthusiastic mission expert focused on education rather than simulation control. Character configuration occurs through `MissionConfig` ScriptableObject assets:

- `ISS_Config.asset`: Character name “Anastasia”, personality “Professional engineer - clear, technical, friendly”
- `Hubble_Config.asset`: Hubble Space Telescope mission specialist
- `Voyager_Config.asset`: Voyager interplanetary mission specialist

The `specialistIntroTemplate` generates concise 40-word, 10-15 second greetings acknowledging the routing context from `route_to_mission`.

A.2 Tool Schema and Validation

The eight tools are defined in `ToolSchemas.json` (169 lines) with complete JSON Schema specifications. Table A.2 documents parameter constraints enforced by the Tool-Registry validation system.

TABLE A.2 – Tool Parameter Constraints

Tool	Parameter	Constraint
<code>create_circular_orbit</code>	<code>altitude_km</code>	160–35,786 km
<code>create_circular_orbit</code>	<code>inclination_deg</code>	0–180°
<code>create_elliptical_orbit</code>	<code>periapsis_km</code>	160–35,786 km
<code>create_elliptical_orbit</code>	<code>apoapsis_km</code>	160–100,000 km
<code>create_elliptical_orbit</code>	<code>inclination_deg</code>	0–180°
<code>set_simulation_speed</code>	<code>speed_multiplier</code>	0.1–100×

A.2.1 Tool Execution Pipeline

The `ToolExecutor` class receives validated tool calls from the agent system and invokes corresponding C# methods:

- **Orbit Tools** → `OrbitController.CreateCircularOrbit()`, `CreateEllipticalOrbit()`
- **Time Controls** → `TimeController.SetSpeed()`, `Pause()`, `ResetTime()`
- **Navigation Tools** → `SceneTransitionManager.TransitionToMission()`, `TransitionToHub()`
- **Workspace** → `OrbitController.ClearOrbit()`

Execution results—success status, generated orbital parameters, error messages—feed back into the LLM response generation cycle through the response prompt template.

A.3 Conversation Context Management

The `ConversationHistory` class maintains conversation continuity across multi-turn dialogues and scene transitions. Table A.3 documents the exchange data structure.

TABLE A.3 – Conversation Exchange Data Structure

Field	Content
timestamp	DateTime of exchange
userMessage	User’s natural language input
agentResponse	Agent’s generated response
toolExecuted	Tool name (or null if conversational)
location	Current scene (Hub, ISS, Hubble, Voyager)

A.3.1 Context Window Management

The system maintains a sliding window of the last 10 exchanges (`maxHistorySize = 10`). Two methods provide context injection into prompts:

- `GetFormattedHistory(lastNExchanges = 5)`: Returns detailed history with timestamps, locations, and tool executions for the last 5 exchanges
- `GetContextSummary(lastNExchanges = 3)`: Returns condensed 3-exchange summary optimized for token efficiency

A.3.2 Cross-Scene Persistence

Scene transitions preserve conversation history through Unity’s `DontDestroyOnLoad` mechanism. The `PromptConsole` GameObject, containing the `ConversationHistory` component, persists across scene unloading when users invoke `route_to_mission` or `return_to_hub` tools. This ensures unbroken dialogue continuity: a user can ask “What was the ISS orbit altitude I created in the Hub?” after transitioning to the ISS Mission Space.

A.4 API Integration

The `OpenAIClient` class (150 lines) implements asynchronous HTTP communication with OpenAI’s Responses API endpoint (<https://api.openai.com/v1/responses>).

A.4.1 Request Structure

Requests to the `/responses` endpoint include:

- **Model:** "gpt-4.1"
- **Input:** User’s natural language message

- **Instructions:** Concatenated system prompt + conversation history + tool schemas

The `CompleteAsync()` method constructs JSON payloads using Unity's `UnityWebRequest` for `async/await` compatibility.

A.4.2 Response Parsing

The client extracts assistant text from JSON responses through a two-stage fallback:

1. **Primary:** Extract `output_text` convenience field (if present)
2. **Fallback:** Concatenate all `output[] .content[] .text` arrays

Tool call JSON undergoes validation by `ToolRegistry` before execution. Results format back into natural language through the response prompt template system, generating contextual explanations like: “I’ve created a circular orbit at 420 km altitude with 51.6° inclination. The orbital velocity is 7.66 km/s, matching the ISS configuration.”

A.4.3 Mission-Specific Configuration

Each Mission Space scene loads scene-specific `OpenAISettings` ScriptableObject assets that override the default system prompt, enabling character switching when users transition from Hub (Mission Control) to Mission Spaces (specialist agents).

Appendix B - Orbital Physics Implementation

This appendix provides detailed technical specifications for the orbital physics simulation engine described in Section 3.7.2, implementing the two-body Keplerian mechanics and vis-viva equation principles outlined in Section 2.3. All equations, algorithms, class methods, and numerical values are verified against the Unity project physics implementation.

B.1 Two-Body Keplerian Mechanics

The simulation implements two-body orbital mechanics under the following simplifying assumptions:

- Earth modeled as a point mass at the coordinate system origin
- Satellite treated as a massless test particle (no gravitational influence on Earth)
- No atmospheric drag, solar radiation pressure, or third-body perturbations
- Instantaneous orbital maneuvers (no finite burn durations)

These assumptions yield closed-form Keplerian solutions suitable for educational visualization while maintaining physical accuracy for the mission profiles studied (ISS, Hubble, Voyager departure trajectory).

B.2 Vis-Viva Equation Implementation

The vis-viva equation relates orbital velocity to position and total orbital energy. Table B.1 documents the physical constants used throughout the simulation.

TABLE B.1 – Physical Constants for Orbital Calculations

Constant	Value	Symbol
Earth's standard gravitational parameter	398,600 km ³ /s ²	μ
Earth's mean radius	6,371 km	R_{\oplus}
Unity scale compression factor	0.000785 Unity/km	k
Unity Earth radius	5 Unity units	R_{Unity}

B.2.1 Circular Orbit Calculation

Circular orbits ($e = 0$) simplify the vis-viva equation to:

$$v_{\text{circular}} = \sqrt{\frac{\mu}{r}} \quad (\text{B.1})$$

where $r = R_{\oplus} + h$ is the orbital radius from Earth's center, and h is the altitude above Earth's surface.

Implementation Method

The `OrbitController.CreateCircularOrbit()` method (lines 229–290) accepts altitude in kilometers and automatically calculates orbital velocity, eliminating user confusion between altitude and speed parameters. Algorithm B.2 documents the calculation sequence.

TABLE B.2 – Circular Orbit Calculation Algorithm

Step	Calculation
1. Validate altitude	$h_{\text{input}} \rightarrow \text{Clamp}(160, 35, 786)$ km
2. Compute orbital radius	$r = R_{\oplus} + h = 6,371 + h$ km
3. Calculate orbital velocity	$v = \sqrt{\mu/r}$ km/s
4. Convert to Unity scale	$r_{\text{Unity}} = R_{\text{Unity}} + h \cdot k$
5. Convert to angular velocity	$\omega = (v \cdot k)/r_{\text{Unity}}$ rad/s

Example: ISS Orbital Velocity

For the International Space Station at $h = 420$ km altitude:

$$r = 6,371 + 420 = 6,791 \text{ km} \quad (\text{B.2})$$

$$v = \sqrt{\frac{398,600}{6,791}} = 7.66 \text{ km/s} \quad (\text{B.3})$$

$$\omega = \frac{7.66 \times 0.000785}{5 + (420 \times 0.000785)} = 0.00113 \text{ rad/s} \quad (\text{B.4})$$

This matches the real ISS orbital velocity of approximately 7.66 km/s.

B.2.2 Elliptical Orbit Calculation

Elliptical orbits ($0 < e < 1$) use the full vis-viva equation:

$$v = \sqrt{\mu \left(\frac{2}{r} - \frac{1}{a} \right)} \quad (\text{B.5})$$

where a is the semi-major axis and r is the instantaneous distance from Earth's center.

Orbital Elements Derivation

Given periapsis altitude h_p and apoapsis altitude h_a :

$$r_p = R_{\oplus} + h_p \quad (\text{periapsis radius}) \quad (\text{B.6})$$

$$r_a = R_{\oplus} + h_a \quad (\text{apoapsis radius}) \quad (\text{B.7})$$

$$a = \frac{r_p + r_a}{2} \quad (\text{semi-major axis}) \quad (\text{B.8})$$

$$e = \frac{r_a - r_p}{r_a + r_p} \quad (\text{eccentricity}) \quad (\text{B.9})$$

Implementation Method

The `OrbitController.CreateEllipticalOrbit()` method (lines 300–367) computes velocity at periapsis using Equation B.5 with $r = r_p$:

$$v_p = \sqrt{\mu \left(\frac{2}{r_p} - \frac{1}{a} \right)} \quad (\text{B.10})$$

Table B.3 documents parameter validation constraints enforced before calculation.

TABLE B.3 – Elliptical Orbit Parameter Constraints

Parameter	Constraint
Periapsis altitude h_p	160–35,786 km
Apoapsis altitude h_a	$h_p + 1$ km to 100,000 km
Eccentricity e	$0 < e < 1$ (enforced implicitly)
Inclination i	0–180°

B.3 Scale Compression

The simulation implements logarithmic scale compression to fit orbital mechanics within the Meta Quest 3’s comfortable rendering volume while preserving geometric relationships.

B.3.1 Compression Factor Derivation

Earth’s physical radius (6,371 km) maps to 5 Unity units:

$$k = \frac{R_{\text{Unity}}}{R_{\oplus}} = \frac{5}{6,371} = 0.000785 \text{ Unity units/km} \quad (\text{B.11})$$

Example Mappings

Table B.4 shows real-world altitudes mapped to Unity rendering coordinates.

TABLE B.4 – Scale Compression Examples

Mission	Real Altitude (km)	Unity Altitude
ISS	420	$420 \times 0.000785 = 0.33$
Hubble	540	$540 \times 0.000785 = 0.42$
Geostationary	35,786	$35,786 \times 0.000785 = 28.1$

This compression maintains visual proportions: the ISS appears at $0.33/5 = 6.6\%$ of Earth’s radius above the surface, matching the real ratio of $420/6,371 = 6.6\%$.

B.3.2 Numerical Stability

All physics calculations occur in real units (km, km/s) before conversion to Unity space for rendering. This ensures:

- No floating-point precision loss from working with very small Unity coordinates
- Physical accuracy verifiable against published orbital data

- Separation of physics (model) from rendering (view)

The `OrbitController` methods perform calculations in kilometers, then convert final results through multiplication by k only when setting Unity Transform positions.

B.4 Trajectory Visualization

The `OrbitVisualizer` class (280 lines) generates trajectory curves by sampling the orbital ellipse equation at discrete points and rendering through Unity's LineRenderer system.

B.4.1 Orbital Ellipse Equation

The orbit trajectory follows the polar equation:

$$r(\theta) = \frac{a(1 - e^2)}{1 + e \cos \theta} \quad (\text{B.12})$$

where θ is the true anomaly (angle from periapsis), a is the semi-major axis, and e is the eccentricity.

B.4.2 Sampling Algorithm

The `CalculateOrbitalPoint()` method (lines 217–231) samples Equation B.12 at 128 evenly-spaced true anomaly angles $\theta \in [0, 2\pi]$. For each sample point:

$$r = \frac{a(1 - e^2)}{1 + e \cos \theta} \quad (\text{B.13})$$

$$x = r \cos(\theta + \omega) \quad (\text{B.14})$$

$$z = r \sin(\theta + \omega) \quad (\text{B.15})$$

where ω is the argument of periapsis (orientation of the ellipse major axis in the orbital plane).

Inclination Transformation

The resulting planar coordinates undergo rotation by inclination angle i via rotation matrix:

$$\mathbf{R}_i = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos i & -\sin i \\ 0 & \sin i & \cos i \end{bmatrix} \quad (\text{B.16})$$

This tilts the orbital plane from equatorial (XZ) to the specified inclination angle, enabling visualization of polar orbits (ISS at 51.6°) and equatorial orbits (geostationary at 0°).

B.4.3 Rendering Configuration

Table B.5 documents the Unity LineRenderer configuration for optimal VR visibility.

TABLE B.5 – LineRenderer Configuration Parameters

Parameter	Value
Path resolution	128 points
Line width	0.05 Unity units
Color	Cyan (0, 1, 1) with 0.7 alpha
Shader	Sprites/Default (view-aligned billboard)
Loop closure	Enabled (connects point 127 to point 0)

Special Cases

- **Circular orbits** ($e = 0$): Simplify to constant radius $r = a$, producing perfect circles
- **Elliptical orbits** ($0 < e < 1$): Render with visible eccentricity
- **Debug visualization:** Green gizmo at periapsis, red gizmo at apoapsis for development testing

B.5 Coordinate System Conventions

The simulation uses Unity's left-handed coordinate system with the following conventions:

- **Origin:** Earth's center of mass
- **Equatorial plane:** XZ plane ($y = 0$)
- **Polar axis:** +Y direction points toward North Pole
- **Reference direction:** +X axis defines 0° longitude

- **Orbital motion:** Counterclockwise when viewed from above North Pole (right-hand rule)

This convention aligns with standard aerospace engineering practices while accommodating Unity's left-handed rendering system.

Appendix C - Voice Pipeline Implementation

This appendix provides detailed technical specifications for the bidirectional voice system described in Section 3.7.3, implementing voice interaction as a hands-free modality for immersive VR environments outlined in Section 2.1. All API endpoints, audio formats, class methods, and processing parameters are verified against the Unity project voice integration code.

C.1 System Architecture

The voice pipeline implements bidirectional audio through ElevenLabs cloud APIs, enabling natural spoken interaction with the agent system. Figure ?? conceptually illustrates the data flow (implementation in code).

Component Responsibilities

- **PromptConsole:** Manages microphone capture, push-to-talk input detection, and audio playback
- **ElevenLabsClient:** Handles HTTP communication with ElevenLabs APIs (381 lines)
- **Unity AudioSource:** Plays synthesized speech through Quest 3's spatial audio system
- **Unity Microphone:** Captures user voice input at 16 kHz sample rate

C.2 Speech-to-Text Pipeline

Speech recognition converts user voice input to text through ElevenLabs' Scribe v1 transcription model. Table C.1 documents the audio capture specifications.

TABLE C.1 – Speech-to-Text Audio Capture Specifications

Parameter	Value
Sample rate	16,000 Hz (optimized for speech)
Bit depth	16-bit PCM
Channels	Mono
Maximum duration	30 seconds
Audio format (transmitted)	WAV with RIFF header
API endpoint	/speech-to-text
Model	scribe_v1

C.2.1 Push-to-Talk Input Detection

Voice recording activates through push-to-talk button press. The `PromptConsole.Update()` method (lines 281–323) implements platform-specific input detection with debouncing to prevent accidental double-triggers.

Input Source Detection

- **Desktop testing:** Space key via `Input.GetKeyDown(KeyCode.Space)`
- **VR deployment:** Quest 3 right controller A button via `OVRInput.Get(OVRInput.Button.One, OVRInput.Controller.RTouch)`

State Machine

Table C.2 documents the recording state transitions.

TABLE C.2 – Recording State Machine

State	Trigger	Next State
Idle	Button press	Recording
Recording	Button release	Processing
Processing	Transcription complete	Idle

During the Recording state, a red visual indicator displays “Listening...” to provide user feedback.

C.2.2 Audio Capture and Conversion

The `StartRecording()` method initiates Unity’s `Microphone.Start()` with the specifications in Table C.1. When the user releases the button, `StopRecordingAndTranscribe()` processes the captured audio.

WAV Conversion Algorithm

The `ConvertAudioClipToWav()` method (lines 322–368) converts Unity’s `AudioClip` format to WAV for API transmission:

1. Extract float samples from `AudioClip.GetData()`
2. Convert float [-1.0, 1.0] to 16-bit signed integer [-32768, 32767]
3. Construct RIFF WAV header (44 bytes):
 - Chunk ID: “RIFF”
 - Format: “WAVE”
 - Subchunk 1: “fmt ” (audio format specification)
 - Subchunk 2: “data” (PCM samples)
4. Concatenate header + PCM data

C.2.3 API Request Structure

The WAV bytes transmit to ElevenLabs via `WWWForm` multipart HTTP POST:

```
POST https://api.elevenlabs.io/v1/speech-to-text
Content-Type: multipart/form-data
Headers: xi-api-key: [API_KEY]
```

Body:

- `file: recording.wav` (binary WAV data)
- `model_id: "scribe_v1"`

C.2.4 Response Parsing

The API returns JSON containing:

- **text**: Transcribed text string
- **confidence**: Recognition confidence score [0.0–1.0]

The transcribed text feeds directly into the agent’s `ProcessUserInput()` method for intent interpretation and tool selection.

C.3 Text-to-Speech Pipeline

Agent text responses convert to speech through ElevenLabs' text-to-speech API. Table C.3 documents the synthesis configuration.

TABLE C.3 – Text-to-Speech Synthesis Parameters

Parameter	Value
Model	<code>eleven_flash_v2_5</code>
Stability	0.7 (voice consistency)
Similarity boost	0.8 (voice clarity)
Speed	1.0 (normal playback)
Audio format (received)	MP3
API endpoint	<code>/text-to-speech/{voiceId}</code>
Synthesis latency	1–3 seconds (typical)

C.3.1 API Request Structure

The `TextToSpeechAsync()` method (lines 31–129) sends synthesis requests:

```
POST https://api.elevenlabs.io/v1/text-to-speech/{voiceId}
Content-Type: application/json
Headers: xi-api-key: [API_KEY]
```

Body:

```
{
  "text": "Agent response text here",
  "model_id": "eleven_flash_v2_5",
  "voice_settings": {
    "stability": 0.7,
    "similarity_boost": 0.8,
    "speed": 1.0
  }
}
```

C.3.2 MP3 Decoding and Playback

The API returns MP3-encoded audio via HTTP response body. The `ConvertMp3ToAudioClipAsync` method (lines 135–154) performs decoding:

1. Write MP3 bytes to temporary file in `Application.temporaryCachePath`

2. Load via `UnityWebRequestMultimedia.GetAudioClip(uri, AudioType.MPEG)`
3. Enable streaming mode for memory efficiency
4. Extract `AudioClip` from request
5. Delete temporary file in `finally` block

The resulting `AudioClip` plays through Unity's `AudioSource` component attached to the camera, utilizing Quest 3's spatial audio capabilities for immersive voice delivery positioned at the user's head location.

C.3.3 Model Selection Rationale

The `eleven_flash_v2_5` model balances:

- **Synthesis speed:** 1–3 seconds for typical 2–3 sentence responses (critical for real-time interaction)
- **Voice fidelity:** Natural prosody and intonation
- **API cost:** Flash models optimize for speed over maximum quality

C.4 Character Voice Management

Each agent character uses a distinct ElevenLabs voice ID configured in `MissionConfig.specialistVoice` ScriptableObject references. Table C.4 documents character voice assignments.

TABLE C.4 – Character Voice ID Assignments

Character	Voice ID	Characteristics
Mission Control	NOpBlnGIn09m6vDvFkFC	Authoritative, encouraging, professional
Anastasia (ISS Specialist)	ZF6FPAbjXT4488VcRRnw	Professional engineer - clear, technical, friendly
Dr. Harrison (Hubble Specialist)	M4zkunnpRihDKTNF0D7f	Veteran aerospace engineer - technical, experienced, proud of Hubble's legacy
Karl (Voyager Specialist)	t1oG321G6Z6edP2XJLiz	Philosophical scientist and cosmic poet - contemplative, poetic, awe-inspiring

C.4.1 Scene-Specific Voice Switching

When users invoke the `route_to_mission` tool, the scene transition loads mission-specific `ElevenLabsSettings` assets that override the default voice ID. This ensures:

- Hub agent responses use Mission Control voice
- ISS Mission Space responses use Anastasia's voice profile
- Each specialist maintains consistent vocal identity

Voice synthesis parameters (stability, similarity boost, speed) remain constant across all characters to maintain audio quality consistency, while the underlying voice models provide tonal and character differentiation.

C.4.2 Voice Settings Persistence

The `ElevenLabsClient` caches the current `ElevenLabsSettings` reference. Scene transitions update this reference automatically through Unity's scene loading hooks, enabling seamless character voice switching without code changes in the agent logic.

C.5 Error Handling and Fallbacks

The voice pipeline implements robust error handling for network failures and API timeouts:

Speech-to-Text Errors

- **Microphone unavailable:** Display error message, fall back to text input
- **API timeout:** Retry once with exponential backoff, then show error
- **Low confidence score:** Accept transcription but log warning

Text-to-Speech Errors

- **API timeout:** Display text response without audio
- **MP3 decode failure:** Log error, display text fallback
- **AudioSource unavailable:** Silent failure, text remains visible

All errors log to Unity console with structured error messages for debugging while maintaining graceful degradation of user experience.

C.6 Performance Optimization

C.6.1 Memory Management

- Temporary WAV/MP3 files deleted immediately after use
- AudioClip instances released when playback completes
- Streaming mode for MP3 decoding reduces peak memory usage
- No audio caching (prioritizes memory over latency)

C.6.2 Latency Budget

Table C.5 documents typical latency components for the complete voice interaction cycle.

TABLE C.5 – Voice Interaction Latency Budget

Component	Latency
User speech duration	Variable (user-controlled)
WAV conversion	< 100 ms
STT API request	500–1500 ms
Agent processing (GPT-4.1)	1000–3000 ms
TTS API request	1000–3000 ms
MP3 decode	< 200 ms
Audio playback start	< 50 ms
Total (excluding user speech)	2.5–7.8 seconds

The 2.5–7.8 second response time falls within acceptable bounds for educational conversational interfaces, where thoughtful responses outweigh instantaneous feedback.

Appendix D - VR Deployment Configuration

This appendix provides detailed technical specifications for the Meta Quest 3 virtual reality deployment described in Section 3.7.4, implementing the spatial learning and immersive presence principles outlined in Section 2.1. All build settings, input mappings, rendering configurations, and scene architecture details are verified against the Unity project configuration files.

D.1 Quest 3 Android Build Configuration

The application deploys to Meta Quest 3 through Unity’s Android build pipeline with OpenXR integration. Table D.1 documents the core build settings from `ProjectSettings.asset`.

TABLE D.1 – Android Build Configuration

Setting	Value
Minimum SDK Version	32 (Android 12L)
Target SDK Version	32 (Android 12L)
Target Architecture	ARMv7 (value: 2)
Graphics API	OpenGL ES 3.0
XR Plugin	OVRPlugin (Oculus SDK)
Stereo Rendering Mode	Single Pass Instanced (value: 2)
Target Device	Meta Quest 3

D.1.1 SDK Version Rationale

Android API level 32 (Android 12L) enables:

- Quest 3’s inside-out tracking system (6DOF head and controller tracking)
- Oculus runtime features (Guardian boundary, passthrough API access)

- Hand tracking capabilities (though not actively used in this application)
- Performance optimizations for Snapdragon XR2 Gen 2 processor

D.1.2 Stereo Rendering Pipeline

Single-pass instanced rendering (value 2 in `ProjectSettings.asset` line 49) reduces CPU overhead by rendering both eye views in a single draw call. This technique:

- Halves per-frame CPU work compared to multi-pass rendering
- Maintains Quest 3's 90 Hz refresh rate target
- Reduces GPU state changes and draw call overhead
- Critical for mobile VR performance on battery-powered hardware

D.1.3 Build Index Scene Configuration

Table D.2 documents the scene inclusion from `EditorBuildSettings.asset`.

TABLE D.2 – Scene Build Index Configuration

Index	Scene Name	File Size
0	Hub.unity	85 KB
1	ISS.unity	65 KB
2	Hubble.unity	66 KB
3	Voyager.unity	62 KB
4	ARHub.unity (experimental)	61 KB

Scene index 0 (Hub) loads at application startup. Scene transitions occur through `SceneManager.LoadSceneAsync()` with scene names or indices.

D.2 Input System Implementation

Controller input integrates Oculus Touch controllers through the OVR Input API. Table D.3 documents the input bindings used in the application.

D.2.1 Push-to-Talk Implementation

The `PromptConsole.Update()` method (line 283) detects the right controller's A button through OVR Input API:

TABLE D.3 – Controller Input Mapping

Action	Desktop	Quest 3 VR
Push-to-talk (voice)	Space key	Right controller A button
Confirm>Select	Enter key	Right controller trigger
Cancel/Back	Escape key	Left controller B button

```
bool aButtonPressed = OVRInput.Get(
    OVRInput.Button.One,
    OVRInput.Controller.RTouch
);
```

State Debouncing

The system tracks previous button state (`_previousAButtonState`) to detect rising edge transitions, preventing accidental double-triggers from single button presses. This ensures one recording session per button press/release cycle.

D.2.2 Desktop Testing Mode

Desktop mode falls back to keyboard input through Unity's legacy Input system:

```
bool spacePressed = Input.GetKeyDown(KeyCode.Space);
```

This enables development iteration without VR hardware, maintaining identical functionality across desktop testing and Quest 3 deployment.

D.2.3 VR Mode Detection

The `StaticVRCameraAligner` class (89 lines) detects VR mode at startup through:

```
bool isVR = XRSettings.isDeviceActive;
```

When `isDeviceActive` returns `true`, the system locates the `OVR Camera Rig` component via `FindObjectOfType<OVR Camera Rig>()` and configures VR-specific camera settings.

D.3 Camera and Rendering Configuration

D.3.1 OVRCameraRig Structure

The Quest 3 camera system follows Oculus SDK conventions. Table D.4 documents the camera hierarchy.

TABLE D.4 – VR Camera Hierarchy

GameObject	Purpose
OVRCameraRig	Root container for VR camera system
TrackingSpace	Offset container for room-scale tracking
CenterEyeAnchor	Head-tracked camera position (stereo)
LeftEyeAnchor	Left eye render camera
RightEyeAnchor	Right eye render camera
LeftHandAnchor	Left controller tracking
RightHandAnchor	Right controller tracking

D.3.2 Near Clip Plane Configuration

The `StaticVRCameraAligner` configures the near clip plane to prevent geometry clipping at close range (line 68):

```
cam.nearClipPlane = 0.01f; // Unity units
```

This 0.01 Unity unit near clip (approximately 1.27 cm in physical space with scale compression factor $k = 0.000785$) ensures UI elements positioned within arm's reach remain visible without clipping.

D.3.3 Desktop Camera Alignment

Desktop mode aligns the fallback camera to match VR positioning conventions, ensuring consistent coordinate systems between development and deployment environments. This allows testing of UI positioning and scene layout without VR hardware.

D.4 Spatial UI Implementation

User interface elements render in 3D world space rather than screen overlay to ensure VR readability and depth perception. Table D.5 documents the UI rendering configuration.

TABLE D.5 – Spatial UI Rendering Configuration

Parameter	Value
Canvas render mode	WorldSpace
Canvas size (transition overlay)	2m × 2m
Canvas distance from camera	1 meter (dynamic)
Mission logo size	512×512 pixels
Text component	TextMeshPro
Background opacity (UI panels)	0.7 alpha
Sort order (transition canvas)	100 (renders on top)

D.4.1 MissionClockUI Pattern

The `MissionClockUI` class (74 lines) demonstrates the spatial UI pattern:

1. Canvas component with `RenderMode.WorldSpace`
2. TextMeshPro text field positioned in 3D environment
3. CanvasGroup component (line 37) controls opacity without render-to-texture overhead
4. Displays mission elapsed time and simulation speed multiplier

D.4.2 Transition Overlay System

The `SceneTransitionManager.CreateTransitionUIIfNeeded()` method (lines 729–838) constructs the transition overlay procedurally:

Canvas Construction

```
Canvas canvas = canvasObj.AddComponent<Canvas>();
canvas.renderMode = RenderMode.WorldSpace;
canvas.sortingOrder = 100;
```

```
RectTransform canvasRect = canvasObj.GetComponent<RectTransform>();
canvasRect.sizeDelta = new Vector2(2f, 2f); // 2m × 2m
```

Dynamic Positioning

The `LateUpdate()` method (lines 154–180) repositions the canvas 1 meter in front of the camera each frame:

```

transitionCanvasTransform.position =
    cachedCameraAnchor.position +
    cachedCameraAnchor.forward * 1f;

transitionCanvasTransform.rotation =
    cachedCameraAnchor.rotation;

```

This dynamic positioning ensures the overlay remains visible during scene transitions when camera references change, avoiding parenting to scene-specific GameObjects that would be destroyed during `SceneManager.LoadSceneAsync()`.

D.5 Scene Architecture and Persistence

The application comprises four navigable scenes sharing common systems through persistent singletons. Table D.6 documents shared components across all scenes.

TABLE D.6 – Common Scene Components

Component	Purpose
OVRCameraRig prefab	VR camera and controller tracking
PromptConsole GameObject	Conversational UI and voice input
TimeController	Simulation speed management
OrbitController	Orbital physics (Hub only)
OrbitVisualizer	Trajectory rendering (Hub only)

D.5.1 Singleton Persistence Mechanism

The `SceneTransitionManager` enforces singleton persistence through Unity's `DontDestroyOnLoad()` mechanism (line 58):

```

if (Instance == null) {
    Instance = this;
    DontDestroyOnLoad(gameObject);
}

```

This ensures the transition UI and conversation context survive scene unloading. Similarly, `ConversationHistory` persists across transitions, preserving the 10-exchange dialogue window.

D.5.2 Asynchronous Scene Loading

Scene loading occurs through `SceneManager.LoadSceneAsync()` with deferred activation (line 252):

```
AsyncOperation loadOperation =
    SceneManager.LoadSceneAsync(sceneName);
loadOperation.allowSceneActivation = false;

// Load scene in background...

// After 4-second logo animation:
loadOperation.allowSceneActivation = true;
```

This deferred activation prevents jarring scene pops, allowing smooth fade-out → logo display → scene activation → fade-in transitions.

D.6 Performance Optimization

Performance optimization targets Quest 3’s mobile GPU constraints. Table D.7 documents the performance budget.

TABLE D.7 – Performance Targets for 90 Hz VR

Metric	Target
Frame time budget	11.1 ms (90 Hz)
Target resolution (per eye)	1832×1920 pixels
Draw calls (Hub scene)	< 100 per frame
Texture memory budget	< 512 MB
Polygon count (visible)	< 100k triangles

D.6.1 Rendering Optimizations

- **Shared material instances:** Reduce draw calls by batching geometry with identical materials
- **Texture compression:** ASTC 6×6 for UI elements, ASTC 4×4 for environment textures
- **Single-pass instanced stereo:** Halves per-frame CPU work (both eyes in one draw call)

- **Occlusion culling:** Disabled (scenes are spatially compact, overhead exceeds benefit)
- **Dynamic batching:** Enabled for small meshes (< 300 vertices)

D.6.2 Memory Management

- Scene file sizes optimized (62–85 KB per scene)
- Texture atlasing for UI sprites
- Audio streaming for voice synthesis (no large audio caching)
- Persistent GameObjects minimized (only transition manager, conversation history)

D.6.3 Frame Time Breakdown

Table D.8 shows typical frame time allocation in the Hub scene (most complex).

TABLE D.8 – Frame Time Budget Breakdown (Hub Scene)

Component	Time
Physics simulation	1.2 ms
Script execution	2.1 ms
Rendering (both eyes)	5.8 ms
VR compositor overhead	1.5 ms
Buffer margin	0.5 ms
Total	11.1 ms

This allocation maintains the 11.1 ms frame budget required for consistent 90 Hz VR without reprojection artifacts (judder).

Appendix E - Complete Demonstration Transcript

This appendix contains the complete transcript of the demonstration session analyzed in Chapter 4. All dialogue has been transcribed from the video recording with light copy-editing for readability. System messages (scene transitions, time acceleration) are indicated with [SIM] annotations.

Cast / Agent Identities

- **USER** — Platform user (learner perspective)
- **CAPCOM** — Mission Control guide
- **ISS_CREW (ANASTASIA)** — ISS astronaut specialist
- **HUBBLE_ENGINEER (DR_HARRISON)** — Hubble Space Telescope specialist
- **VOYAGER_SPECIALIST (CARL)** — Voyager deep-space mission specialist
- **SIM** — Non-diegetic system messages (routing, time control)

Complete Interaction Transcript

Opening Sequence

NARRATOR (Cutscene): Above us, thousands of satellites trace perfect arcs—testaments to gravity, velocity, and the geometry of motion. You're about to join that tradition. I'll guide the physics; you'll design the orbits. Together, we'll explore how spacecraft navigate the cosmos.

Hub Arrival and Initial Interaction

CAPCOM: Welcome to Mission Control. You're floating above Earth. Use your controllers to drift closer. Circle around and watch the continents slide beneath you, oceans catching light, night falling while dawn breaks on the other side. I'm Capcom. We can design orbits together, build trajectories, see how satellites actually move—or I can route you to three specialists:

- a crew member aboard the ISS who lives this work every day,
- an engineer who built Hubble to see what we couldn't from down here,
- and a space enthusiast who dreams about what Voyager means beyond everything.

When you're ready to speak, press A on your right controller. I'll be listening. Enjoy the view—that blue marble beneath you is home.

USER: Hey, Capcom. I see we have a satellite here. Is this the satellite that we can build?

CAPCOM: [Music] Hello from Mission Control. Yes—this is your simulation satellite. We can build orbits by setting circular or elliptical paths. Tell me either the altitude for circular, or periapsis and apoapsis for elliptical, and I'll get it launched.

ISS Mission Consultation (First Visit)

USER: Cool. Let's start with a circular orbit. I need an altitude—but what's a good altitude? What determines a good altitude?

CAPCOM: [Music] Connecting you to the ISS mission for real-world examples. Stand by.

SIM: [ROUTING → ISS_CREW]

ISS_CREW (ANASTASIA): [Music] Hello from the ISS. For circular LEO orbits, engineers balance drag, launch energy, accessibility, and mission needs. The ISS flies around ~420 km—high enough to limit atmospheric drag so we don't slow down too quickly, but not so high that crew/cargo launches become inefficient. It's also reachable from multiple launch sites. Too low: frequent reboosts. Too high: launches get harder and more expensive.

USER: Would it be absurd to choose the same altitude as the ISS?

ISS_CREW (ANASTASIA): Not absurd at all—quite sensible. ~420 km is a proven compromise. Many Earth-observation and crewed missions use similar altitudes for the same reasons. You'd be following decades of flight heritage.

USER: Is the ISS circular?

ISS_CREW (ANASTASIA): Nearly. The path is close to circular with small variations (drag, gravity field). We keep it tight with periodic reboosts. Near-circular simplifies rendezvous and ops.

USER: What was the altitude again? Please route me back to the hub.

ISS_CREW (ANASTASIA): ~420 km. Handing you back. Safe travels.

SIM: [RETURN → CAPCOM]

Creating Circular Orbit

USER: I'll build a circular orbit at 422 km.

CAPCOM: Copy. Circular orbit set at 422 km, inclination 0° (equatorial). Orbital speed ~7.66 km/s—like the ISS, but over the equator.

USER: Can I choose the speed?

CAPCOM: In this sim, speed is derived from altitude by physics. At 422 km, you need ~7.66 km/s for a stable circular orbit. If you want a different speed, pick a different altitude.

ISS Mission Consultation (Second Visit - Orbital Period)

USER: How long does one lap take? It looks fast, but Earth is massive.

CAPCOM: Routing you to the ISS for orbital period context.

SIM: [ROUTING → ISS_CREW]

ISS_CREW (ANASTASIA): At ~420 km, the ISS completes a lap in ~92.8 minutes—about 15 orbits/day. Higher altitude → slower speed → longer period. That's why we see a sunrise/sunset ~every 45 minutes.

SIM: [RETURN → CAPCOM]

USER: With our current orbit (\sim 92 minutes), can we speed up the simulation so I can see it go around?

CAPCOM: Time acceleration: $10\times$. You'll see the \sim 92-minute trip much faster.

SIM: [TIME_WARP $10\times$]

Hubble Mission Consultation (Elliptical Orbits)

USER: Great. I built a circular orbit. What's an elliptical orbit?

CAPCOM: Perfect question for Hubble—routing.

SIM: [ROUTING \rightarrow HUBBLE_ENGINEER]

HUBBLE_ENGINEER (DR_HARRISON): [Music] Circular means constant distance and speed—like Hubble (\approx 540 km) for steady imaging and predictable servicing. Elliptical stretches the path: closer at periapsis (faster), farther at apoapsis (slower). Some science missions pick elliptical to linger where conditions are best or to avoid interference (e.g., high apogee above the radiation belts).

USER: Why choose elliptical if circular is simpler?

HUBBLE_ENGINEER (DR_HARRISON): To optimize time in a region. Example: Chandra uses a highly elliptical orbit to spend long stretches far from Earth's radiation belts for cleaner X-ray observations. Trade-off: more complex tracking/operations.

SIM: [RETURN \rightarrow CAPCOM]

Creating Elliptical Orbits

USER: Back to Mission Control so I can build one.

CAPCOM: You're back at the Hub.

USER: Please build an elliptical orbit—anything interesting so I can compare.

CAPCOM: Elliptical orbit set: periapsis 400 km, apoapsis 2,000 km, inclination 0° . Watch it speed up near Earth and slow down far away.

USER: Can we run at maximum sim speed?

CAPCOM: Time acceleration: $100\times$. Ellipticity will be very clear.

SIM: [TIME_WARP 100×]

USER: Make it more elliptical; it still looks circular.

CAPCOM: Roger. New orbit: periapsis 200 km, apoapsis 1,000 km, inclination 0°. The speed contrast will be much more dramatic.

USER: Speed to maximum... oh, it already is.

CAPCOM: Confirmed—100× time. Enjoy the fast-forward view.

USER: Wow—huge difference between near and far.

Voyager Mission Consultation (Escape Trajectories)

USER: For deep-space missions, I assume there aren't circular orbits at all. Is everything elliptical?

CAPCOM: Routing to Voyager for escape-trajectory talk.

SIM: [ROUTING → VOYAGER_SPECIALIST]

VOYAGER_SPECIALIST (CARL): Great question. To leave Earth, you don't stay in an Earth orbit. Voyager 1 began in a near-circular parking orbit briefly, then received a final boost to a hyperbolic escape trajectory—no longer bound to Earth. Satellites around Earth use circular or elliptical orbits to remain nearby; escape missions transition from a staging orbit to hyperbolic and head out—past Earth, past the Sun's planets, toward interstellar space.

USER: So it's like breaking free from gravity?

VOYAGER_SPECIALIST (CARL): Exactly. Think of it as climbing out of Earth's gravitational well. Circular and elliptical orbits stay in the well, perpetually falling around Earth. A hyperbolic trajectory has enough energy to climb out completely. Voyager didn't just orbit—it escaped.

USER: How did they actually do it?

VOYAGER_SPECIALIST (CARL): Voyager launched from Earth, first reaching a temporary parking orbit around 160-200 km. Then, at the optimal moment, the upper stage fired—adding enough velocity to exceed escape speed. That burn transitioned the trajectory from elliptical to hyperbolic. Once beyond Earth's sphere of influence, Voyager's path was governed by the Sun's gravity, heading toward Jupiter for the first gravitational assist.

USER: That's incredible. It's really leaving everything behind.

VOYAGER SPECIALIST (CARL): Past Earth, past the Sun's planets, toward interstellar space. Voyager carries humanity's first message to the cosmos. The golden record, the images, the sounds of Earth—all riding on a hyperbolic trajectory that will never return. It's physics, yes, but it's also a statement of who we are and what we dream.

End of Recorded Session

[Demonstration recording concludes. User remains in Voyager Mission Space, continuing exploratory dialogue.]

FOLHA DE REGISTRO DO DOCUMENTO

1. CLASSIFICAÇÃO/TIPO TC	2. DATA 25 de março de 2015	3. DOCUMENTO Nº DCTA/ITA/DM-018/2015	4. Nº DE PÁGINAS 114
-----------------------------	--------------------------------	---	-------------------------

5. TÍTULO E SUBTÍTULO:

Educational Orbit Simulation with Generative AI Agentic Workflow and Virtual Reality Visualisation

6. AUTOR(ES):

Eduardo Moura Zindani

7. INSTITUIÇÃO(ÓES)/ÓRGÃO(S) INTERNO(S)/DIVISÃO(ÓES):

Instituto Tecnológico de Aeronáutica – ITA

8. PALAVRAS-CHAVE SUGERIDAS PELO AUTOR:

AI; VR

9. PALAVRAS-CHAVE RESULTANTES DE INDEXAÇÃO:

AI; VR

10. APRESENTAÇÃO:

Nacional Internacional

ITA, São José dos Campos. Curso de Mestrado. Programa de Pós-Graduação em Engenharia Aeronáutica e Mecânica. Área de Sistemas Aeroespaciais e Mecatrônica. Orientador: Prof. Dr. Adalberto Santos Dupont. Coorientadora: Profª. Drª. Doralice Serra. Defesa em 05/03/2015. Publicada em 25/03/2015.

11. RESUMO:

Métodos educacionais tradicionais frequentemente encontram dificuldades para transmitir conceitos complexos, espaciais e dinâmicos como os da mecânica orbital. A recente convergência entre a Realidade Mista (RM) de consumo e os sofisticados agentes de Inteligência Artificial (IA) Generativa apresenta uma oportunidade para criar um novo paradigma de interfaces de aprendizagem intuitivas e experenciais. Este trabalho detalha o projeto, desenvolvimento e demonstração de uma plataforma educacional interativa para a exploração dos princípios da mecânica orbital. O objetivo principal do sistema é conectar a teoria de leis físicas abstratas à compreensão intuitiva, permitindo que os usuários aprendam por meio da interação corporificada. A metodologia é centrada em uma arquitetura modular que integra dois componentes principais: (1) um "cérebro" agente generativo, impulsionado por Modelos de Linguagem Abrangentes, que interpreta comandos em linguagem natural e atua como um guia educacional especializado; e (2) um "mundo" de simulação em tempo real e visualização imersiva em realidade mista, construído no motor Unity para o Meta Quest 3, que renderiza trajetórias orbitais fisicamente precisas em um espaço tridimensional onde os usuários vivenciam a mecânica orbital de dentro. A plataforma facilita um ciclo de interação multimodal contínuo, onde os comandos de voz do usuário são capturados, processados pelo agente para alterar os parâmetros da simulação e refletidos na visualização imersiva em RV com feedback auditivo conversacional. Este trabalho entrega um protótipo funcional que demonstra uma nova abordagem para a educação científica, transformando dados abstratos em uma experiência manipulável e conversacional para promover uma aprendizagem exploratória e profundamente engajadora. A plataforma é disponibilizada como software de código aberto para permitir validação, adaptação e extensão pela comunidade para diversos contextos educacionais.

12. GRAU DE SIGILO:

OSTENSIVO

RESERVADO

SECRETO